

ECE 550D

Fundamentals of Computer Systems and Engineering

Fall 2023

Instruction Set Architectures (ISAs) and MIPS (Microprocessor without Interlocked Pipeline Stages)

Xin Li & Dawei Liu
Duke Kunshan University

Slides are derived from work by
Andrew Hilton, Tyler Bletsch and Rabi Younes (Duke)

Last time...

- Who can remind us what we did last time?
 - Finite State Machines
 - Division

Now: Moving to software


- C is pretty low level
 - Most of you: in 551, can program in C
 - Others: assume you know programming, pointers, etc.
 - But compiled into **assembly**...
- Lower level programming
 - What does assembly look like?
 - How does software communicate with hardware?
 - What variations can there be?
 - Advantages/disadvantages?

3

Assembly

- Assembly programming:
 - 1 **machine instruction** at a time
 - Still in “human readable form”
 - add \$1, \$2, \$3
 - Much “lower level” than any other programming
 - Limited number of **registers** vs unlimited variables
 - Flat scope

How to say “\$1” out loud:



4

Registers

- Two places processors can store data
 - Registers (saw these---sort of):
 - In processor
 - Few of them (e.g., 32)
 - Fast (more on this much later in semester)
 - Memory (later):
 - Outside of processor
 - Huge (e.g., 16GB)
 - Slow (generally about 100—200x slower than registers, more later)
- For now: think of registers like “variables”
 - But only 32 of them
 - E.g., $\$1 = \$2 + \$3$ much like $x = y + z$

5

Simple, Running Example

```
// silly C code
```

```
int sum, temp, x, y;  
while (true){  
    temp = x + y;  
    sum = sum + temp;  
}
```

```
// equivalent MIPS assembly code
```

```
loop:  lw $1, Memory[1004]  
       lw $2, Memory[1008]  
       add $3, $1, $2  
       add $4, $4, $3  
       j loop
```

*Memory references
don't quite work
like this...we'll
correct this later.*

OK, so what does this assembly code mean?
Let's dig into each line ...

6

Simple, Running Example

```
loop:  lw $1, Memory[1004]
      lw $2, Memory[1008]
      add $3, $1, $2
      add $4, $4, $3
      j loop
```

NOTES

"loop:" = line label (in case we need to refer to this instruction's PC)

lw = "load word" = read a word (32 bits) from memory

\$1 = "register 1" → put result read from memory into register 1

Memory[1004] = address in memory to read from (where x lives)

Note: PC means "program counter"

Note: almost all MIPS instructions put destination (where result gets written) first (in this case, \$1)

7

Simple, Running Example

```
loop:  lw $1, Memory[1004]
      lw $2, Memory[1008]
      add $3, $1, $2
      add $4, $4, $3
      j loop
```

NOTES

lw = "load word" = read a word (32 bits) from memory

\$2 = "register 2" → put result read from memory into register 2

Memory[1008] = address in memory to read from (where y lives)

8

Simple, Running Example

```
loop:  lw $1, Memory[1004]
       lw $2, Memory[1008]
       add $3, $1, $2
       add $4, $4, $3
       j loop
```

NOTES

add \$3, \$1, \$2= add what's in \$1 to what's in \$2 and put result in \$3

9

Simple, Running Example

```
loop:  lw $1, Memory[1004]
       lw $2, Memory[1008]
       add $3, $1, $2
       add $4, $4, $3
       j loop
```

NOTES

add \$4, \$4, \$3= add what's in \$4 to what's in \$3 and put result in \$4

Note: this instruction overwrites previous value in \$4

10

Simple, Running Example

```
loop:  lw $1, Memory[1004]
      lw $2, Memory[1008]
      add $3, $1, $2
      add $4, $4, $3
      j loop
```

NOTES

j = "jump"

loop = PC of instruction at label "loop" (the first lw instruction above)
sets next PC to the address labeled by "loop"

Note: all other instructions in this code set next PC = PC++

11

Assembly too high level for machine

- Human readable is not (easily) machine executable
 - add \$1, \$2, \$3
- **Instructions are numbers too!**
 - Bit fields (like FP numbers)
- Instruction Format
 - Establishes a mapping from "instruction" to binary values
 - Which bit positions correspond to which parts of the instruction (operation, operands, etc.)
- In MIPS, each assembly instruction has a unique 32-bit representation:
 - add \$3, \$2, \$7 \leftrightarrow 00000000010001110001100000100000
 - lw \$8, Mem[1004] \leftrightarrow 10001100000010000000001111101100
- **Assembler** does this translation
 - Humans don't typically need to write raw bits

12

What Must be Specified?

- Instruction “opcode”
 - What should this operation do? (add, subtract,...)
- Location of operands and result
 - Registers (which ones?)
 - Memory (what address?)
 - Immediates (what value?)
- Data type and Size
 - Usually included in opcode
 - E.g., signed vs unsigned int (if it matters)
- What instruction comes next?
 - Sequentially next instruction, or jump elsewhere

13

The ISA

- Instruction Set Architecture (ISA)
 - **Contract between hardware and software**
 - Specifies everything hardware and software need to agree on
 - Instruction encoding and effects
 - Memory structure
 - Etc.
- Many different ISAs
 - x86 and x86_64 (Intel and AMD)
 - POWER (IBM)
 - **MIPS**
 - ARM
 - SPARC (Oracle)

14

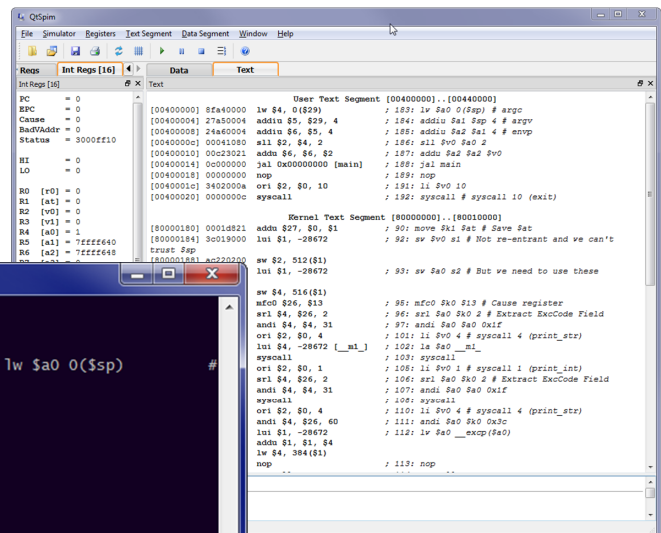
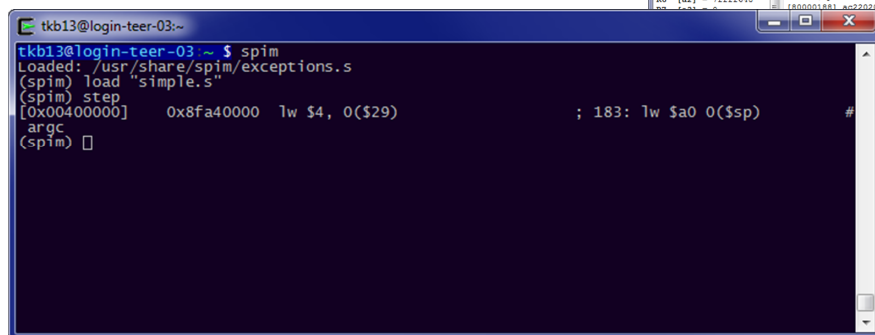
Our focus: MIPS

- We will work with MIPS
 - x86 is ugly (x86_64 is less ugly, but still nasty)
 - MIPS is relatively “clean”
 - More on this in a minute

15

But I don't have a MIPS computer?

- We'll be using SPIM
 - Command line version: spim
 - Graphical version: qtspim
- Edit in emacs, run in SPIM



- Get QtSPIM at: <http://spimsimulator.sourceforge.net/>

16

ISAs: RISC vs CISC

- Two broad categories of ISAs:
 - Complex Instruction Set Computing
 - Came first, days when people always directly wrote assembly
 - Big complex instructions
 - Reduced Instruction Set Computing
 - Goal: make hardware simple and fast
 - Write in high level language, let compiler do the dirty work
 - Rely on compiler to optimize for you
- Note:
 - Sometimes fuzzy: ISAs may have some features of each
 - Common mis-conception: not about how many different insns!

17

ISAs: RISC vs CISC

Reduced Instruction Set Computing

- Simple, fixed length instruction encoding
- Few **memory addressing modes**
- “Many” registers (e.g., 32)
- Three-operand arithmetic (dest = src1 op src2)
- **Load-store** ISA

18

ISAs: RISC vs CISC

- **Complex Instruction Set Computing**
 - Variable length instruction encoding (sometimes quite complex)
 - Many addressing modes, some quite complex
 - Few registers (e.g., 8)
 - Various operand models
 - Stack
 - Two-operand (dest = src op dest)
 - Implicit operands
 - Can operate directly on memory
 - Register = Memory op Register
 - Memory = Memory op Register
 - Memory = Memory op Memory

19

Memory Addressing Modes

- Memory location: how to specify address
 - Simple (RISC)
 - Register + Immediate (e.g., address = \$4 + 16)
 - Register + Register (e.g., address = \$4 + \$7)
 - Complex (CISC)
 - Auto-increment (e.g., address = \$4; \$4 = \$4 + 4;)
 - Scaled Index (e.g., address = \$4 + (\$2 << 2) + 0x1234)
 - Memory indirect (e.g., address = memory[\$4])

20

Load-Store ISA

- Load-store ISA (RISC):
 - Specific instructions (loads/stores) to access memory
 - Loads read memory (and **only** read memory)
 - Stores write memory (and **only** write memory)
- Contrast with (CISC)
 - General memory operands ($\$4 = \text{mem}[\$5] + \$3$)
 - Memory/memory operations: $\text{mem}[\$4] = \text{mem}[\$5] + \$3$


21

Stored Program Computer

- Instructions: a fixed set of built-in operations
- Instructions and data are stored in memory
 - Allows general purpose computation!
- Fetch-Execute Cycle


```
while (!done)
    fetch instruction
    execute instruction
```

- Effectively what hardware does
- This is what the SPIM Simulator does



22

How are Instructions Executed?



- **Instruction Fetch:**
Read instruction bits from memory
- **Decode:**
Figure out what those bits mean
- **Operand Fetch:**
Read registers (+ mem to get sources)
- **Execute:**
Do the actual operation (e.g., add the #s)
- **Result Store:**
Write result to register or memory
- **Next Instruction:**
Figure out mem addr of next insn, repeat

23

More Details on Execution?

- Previous slide provides a high-level overview
 - Called von Neumann model
 - John von Neumann: Eniac (first digital computer)
- More details: How hardware works
 - Later in the course
- Now, diving into assembly programming/MIPS

24

Assembly Programming

- How do you write an assembly program?
- How do you write a program (in general)?

25

5 Step Plan (ECE 551)

- 5 Steps to write any program:
 1. Work an example yourself
 2. Write down what you did
 3. Generalize steps from 2
 4. Test generalized steps on different example
 5. Translate generalized steps to code

26

How to Write a Program

- How I teach programming:

1. Work an example yourself
2. Write down what you did
3. Generalize steps from 2
4. Test generalized steps on different example
5. Translate generalized steps to code


Develop Algorithm
In (Familiar) Higher
Level Language

Then translate to lower level language


27

Why do I bring this up?

- Very Hard:



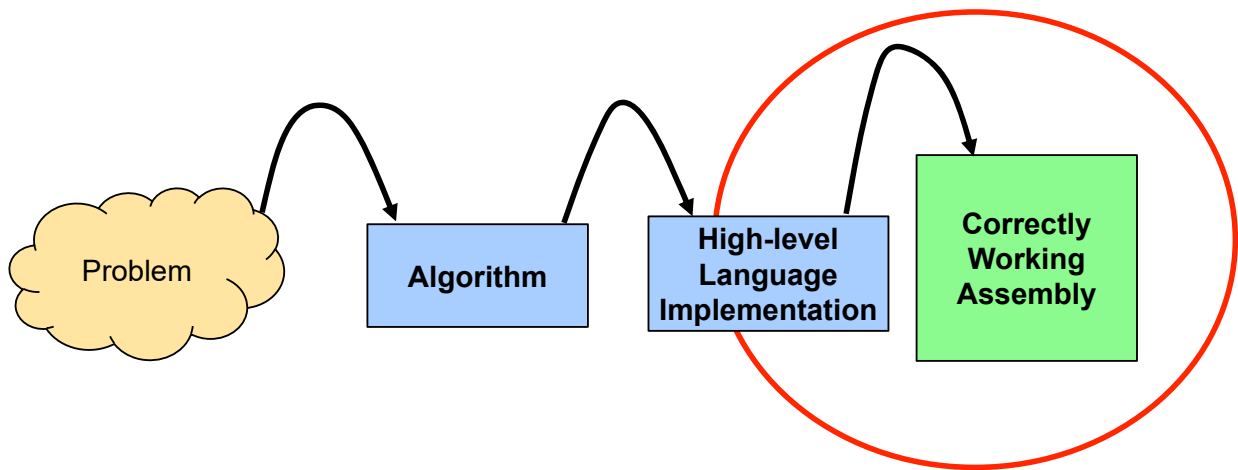
- Easier:



28

Our focus

- We will focus on the assembly step
 - Assume you know how to devise an algorithm for a problem
 - I'll use C.



29

Simplest Operations We Might Want?

- What is the simplest computation we might do?
 - Add two numbers:

$x = a + b;$
add \$1, \$2, \$3

“Add \$2 + \$3, and store it in \$1”

Note: when writing assembly, basically pick reg for a, reg for b, reg for x
Not enough regs for all variables? We'll talk about that later...

30

MIPS Integer Registers

- Recall: registers
 - Fast
 - In CPU
 - Directly compute on them
- 31 x 32-bit GPRs (R0 = 0)
 - GPRs: general purpose registers
- Also floating point registers
- A few special purpose regs too
 - PC = Address of next insn



31

Executing Add

Address	Instruction
1000	add \$8, \$4, \$6
1004	add \$5, \$8, \$7
1008	add \$6, \$6, \$6
100C	...
1010	...

PC tells us where to execute next

Register	Value
\$0	0000 0000
\$1	1234 5678
\$2	C001 D00D
\$3	1BAD F00D
\$4	9999 9999
\$5	ABAB ABAB
\$6	0000 0001
\$7	0000 0002
\$8	0000 0000
\$9	0000 0000
\$10	0000 0000
\$11	0000 0000
\$12	0000 0000
...	...
PC	0000 1000

32

Executing Add

Address	Instruction
1000	add \$8, \$4, \$6
1004	add \$5, \$8, \$7
1008	add \$6, \$6, \$6
100C	...
1010	...

Add reads its source registers, and uses their values directly ("register direct")

9999 9999
0000 0001
 9999 999A

Register	Value
\$0	0000 0000
\$1	1234 5678
\$2	C001 D00D
\$3	1BAD F00D
\$4	9999 9999
\$5	ABAB ABAB
\$6	0000 0001
\$7	0000 0002
\$8	0000 0000
\$9	0000 0000
\$10	0000 0000
\$11	0000 0000
\$12	0000 0000
...	...
PC	0000 1000

33

Executing Add

Address	Instruction
1000	add \$8, \$4, \$6
1004	add \$5, \$8, \$7
1008	add \$6, \$6, \$6
100C	...
1010	...

Add writes its result to its destination register

Register	Value
\$0	0000 0000
\$1	1234 5678
\$2	C001 D00D
\$3	1BAD F00D
\$4	9999 9999
\$5	ABAB ABAB
\$6	0000 0001
\$7	0000 0002
\$8	9999 999A
\$9	0000 0000
\$10	0000 0000
\$11	0000 0000
\$12	0000 0000
...	...
PC	0000 1000

34

Executing Add

Address	Instruction
1000	add \$8, \$4, \$6
1004	add \$5, \$8, \$7
1008	add \$6, \$6, \$6
100C	...
1010	...

And goes to the sequentially next instruction
(PC = PC + 4)

Register	Value
\$0	0000 0000
\$1	1234 5678
\$2	C001 D00D
\$3	1BAD F00D
\$4	9999 9999
\$5	ABAB ABAB
\$6	0000 0001
\$7	0000 0002
\$8	9999 999A
\$9	0000 0000
\$10	0000 0000
\$11	0000 0000
\$12	0000 0000
...	...
PC	0000 1004

35

Executing Add

Address	Instruction
1000	add \$8, \$4, \$6
1004	add \$5, \$8, \$7
1008	add \$6, \$6, \$6
100C	...
1010	...

You all do the next instruction!

Register	Value
\$0	0000 0000
\$1	1234 5678
\$2	C001 D00D
\$3	1BAD F00D
\$4	9999 9999
\$5	ABAB ABAB
\$6	0000 0001
\$7	0000 0002
\$8	9999 999A
\$9	0000 0000
\$10	0000 0000
\$11	0000 0000
\$12	0000 0000
...	...
PC	0000 1004

36

Executing Add

Address	Instruction
1000	add \$8, \$4, \$6
1004	add \$5, \$8, \$7
1008	add \$6, \$6, \$6
100C	...
1010	...



We set \$5 equal to
\$8 (9999 999A) + \$7 (2) = 9999 999C

and PC = PC + 4

Register	Value
\$0	0000 0000
\$1	1234 5678
\$2	C001 D00D
\$3	1BAD F00D
\$4	9999 9999
\$5	9999 999C
\$6	0000 0001
\$7	0000 0002
\$8	9999 999A
\$9	0000 0000
\$10	0000 0000
\$11	0000 0000
\$12	0000 0000
...	...
PC	0000 1008

37

Executing Add

Address	Instruction
1000	add \$8, \$4, \$6
1004	add \$5, \$8, \$7
1008	add \$6, \$6, \$6
100C	...
1010	...



Its perfectly fine to have \$6 as a src and a dst
This is just like $x = x + x$; in C, Java, etc:

$$1 + 1 = 2$$

Register	Value
\$0	0000 0000
\$1	1234 5678
\$2	C001 D00D
\$3	1BAD F00D
\$4	9999 9999
\$5	9999 999C
\$6	0000 0001
\$7	0000 0002
\$8	9999 999A
\$9	0000 0000
\$10	0000 0000
\$11	0000 0000
\$12	0000 0000
...	...
PC	0000 1008

38

Executing Add

Address	Instruction
1000	add \$8, \$4, \$6
1004	add \$5, \$8, \$7
1008	add \$6, \$6, \$6
100C	...
1010	...



Its perfectly fine to have \$6 as a src and a dst
This is just like $x = x + x$; in C, Java, etc:

$$1 + 1 = 2$$

Register	Value
\$0	0000 0000
\$1	1234 5678
\$2	C001 D00D
\$3	1BAD F00D
\$4	9999 9999
\$5	9999 999C
\$6	0000 0002
\$7	0000 0002
\$8	9999 999A
\$9	0000 0000
\$10	0000 0000
\$11	0000 0000
\$12	0000 0000
...	...
PC	0000 100C

39

Summary

- MIPS ISA and Assembly Programming
- Continue on next lecture...

40