

ECE 550D

Fundamentals of Computer Systems and Engineering

Fall 2023

Pipelining

Xin Li & Dawei Liu
Duke Kunshan University

Slides are derived from work by
Andrew Hilton, Tyler Bletsch and Rabih Younes (Duke)

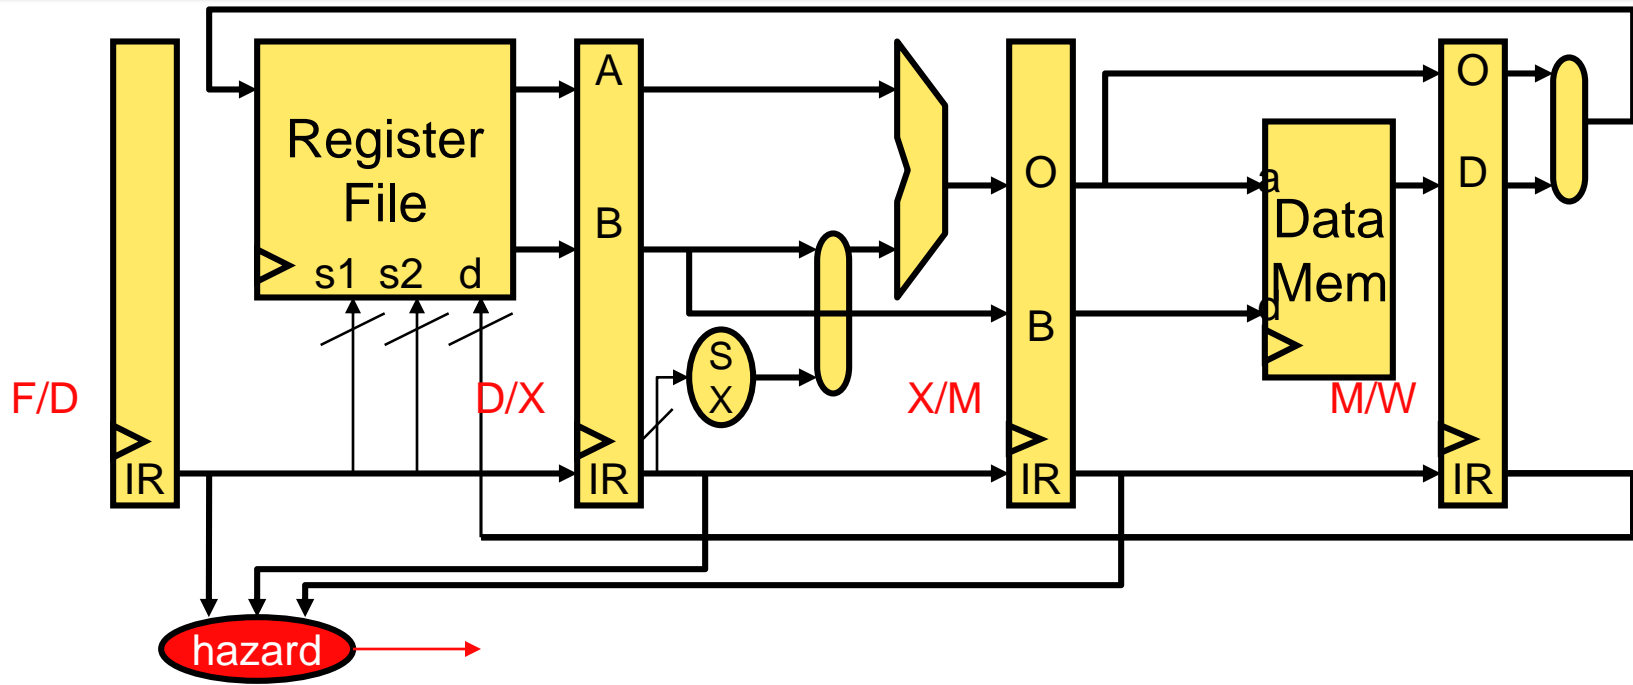
Last Time

- What did we do last time?
- Pipelining
 - Pipelined datapath
 - Pipeline control
 - Dependences and hazards

Hardware Interlocks

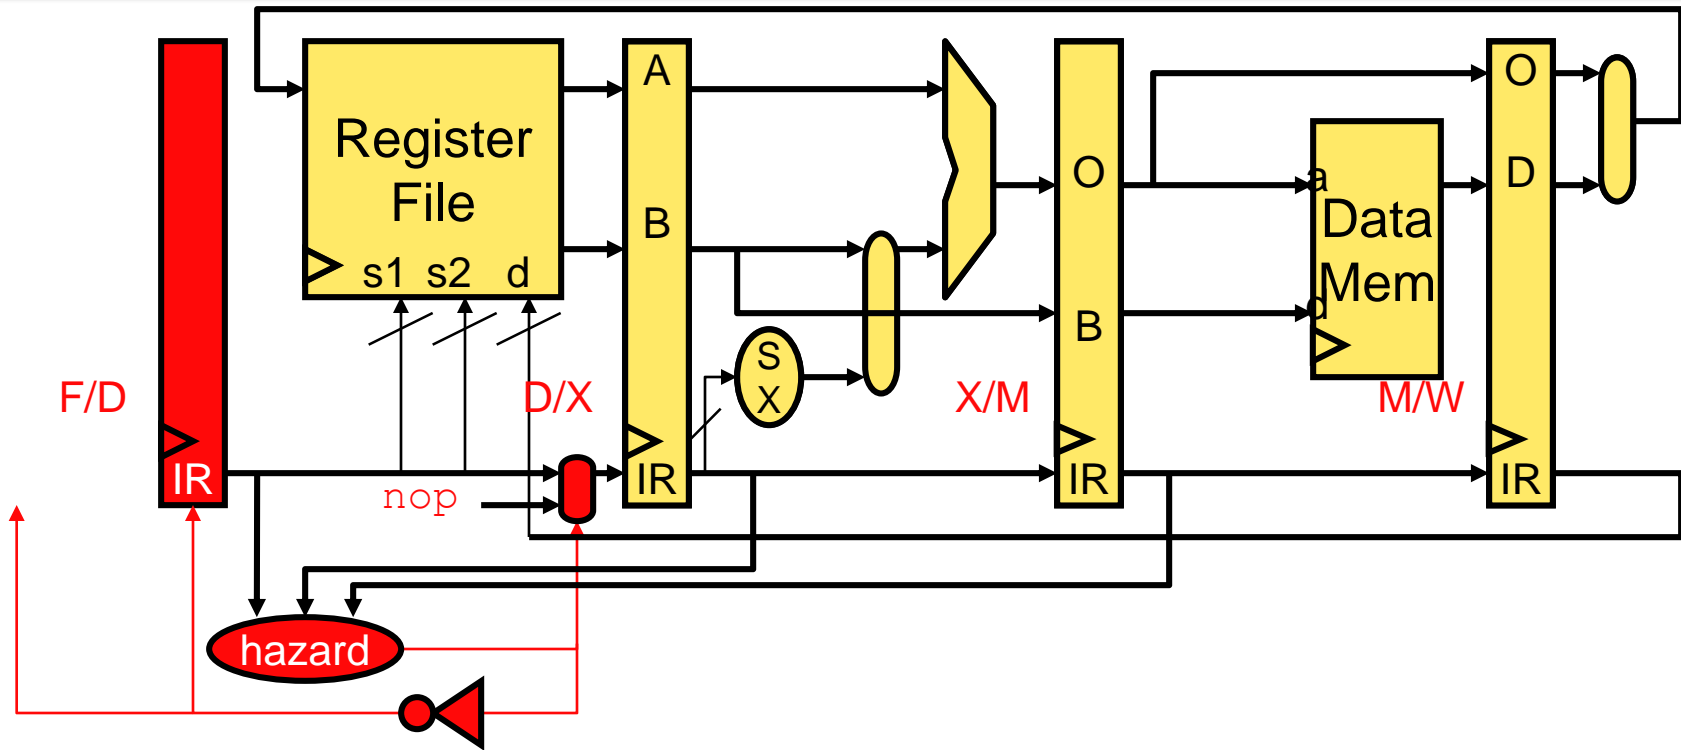
- A better (more compatible) way: **hardware interlocks**
 - Processor detects data hazards and fixes them
 - Two aspects to this
 - Detecting hazards
 - Fixing hazards

Detecting Data Hazards



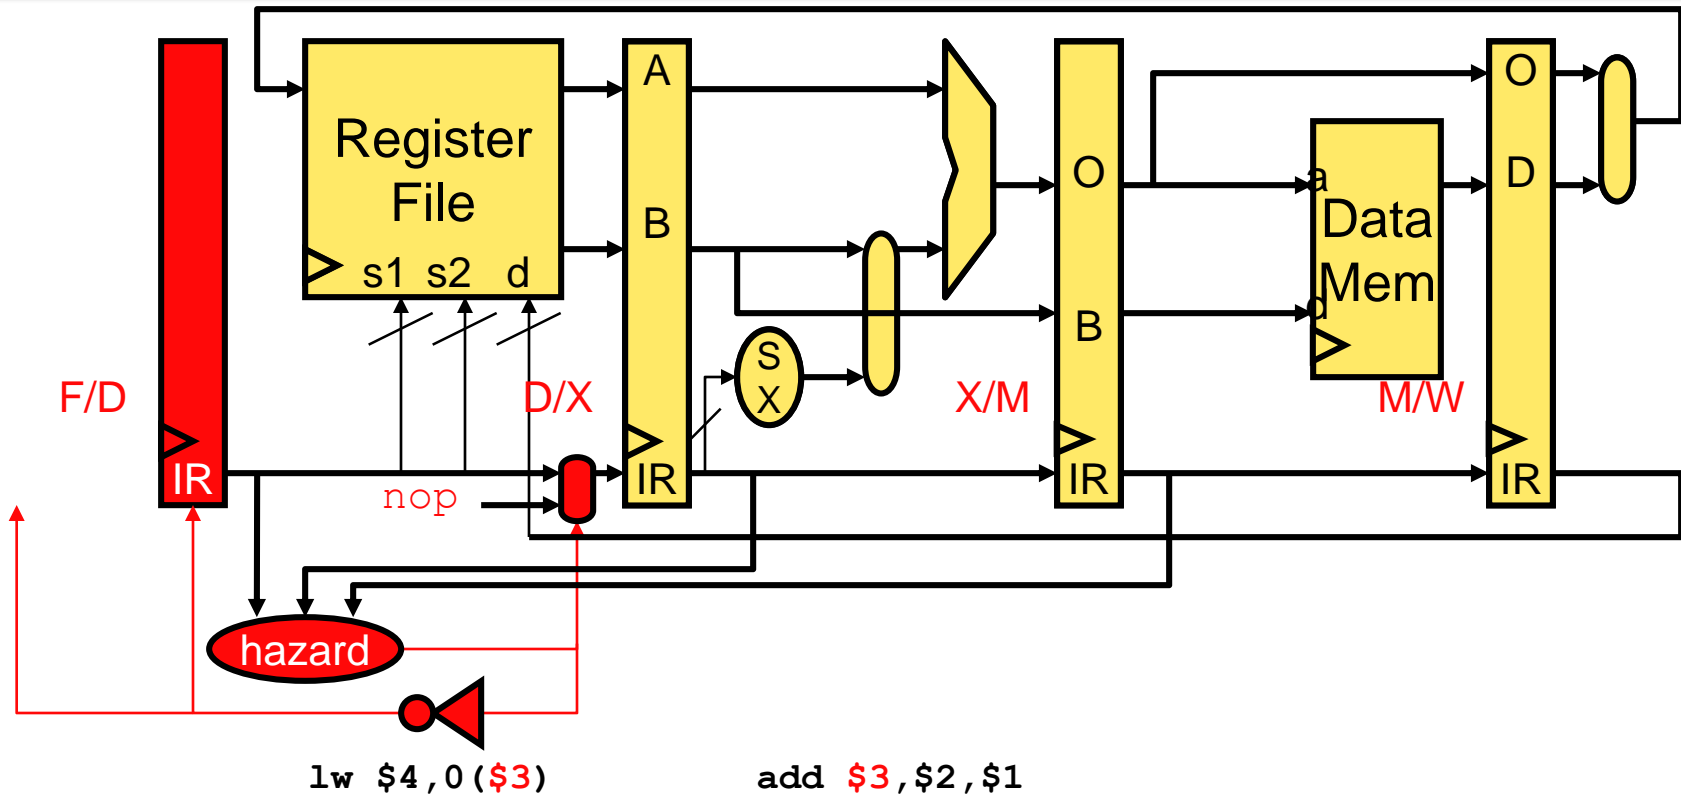
- Compare F/D insn input register names with output register names of older insns in pipeline
 - Hazard =
 - $(F/D.IR.RS1 == D/X.IR.RD) \parallel (F/D.IR.RS2 == D/X.IR.RD) \parallel$
 - $(F/D.IR.RS1 == X/M.IR.RD) \parallel (F/D.IR.RS2 == X/M.IR.RD)$

Fixing Data Hazards



- Prevent F/D insn from reading (advancing) this cycle
 - Write **nop** into D/X.IR (effectively, insert **nop** in hardware)
 - Also reset (clear) the datapath control signals
 - Disable F/D latch and PC write enables (why?)
- Re-evaluate situation next cycle

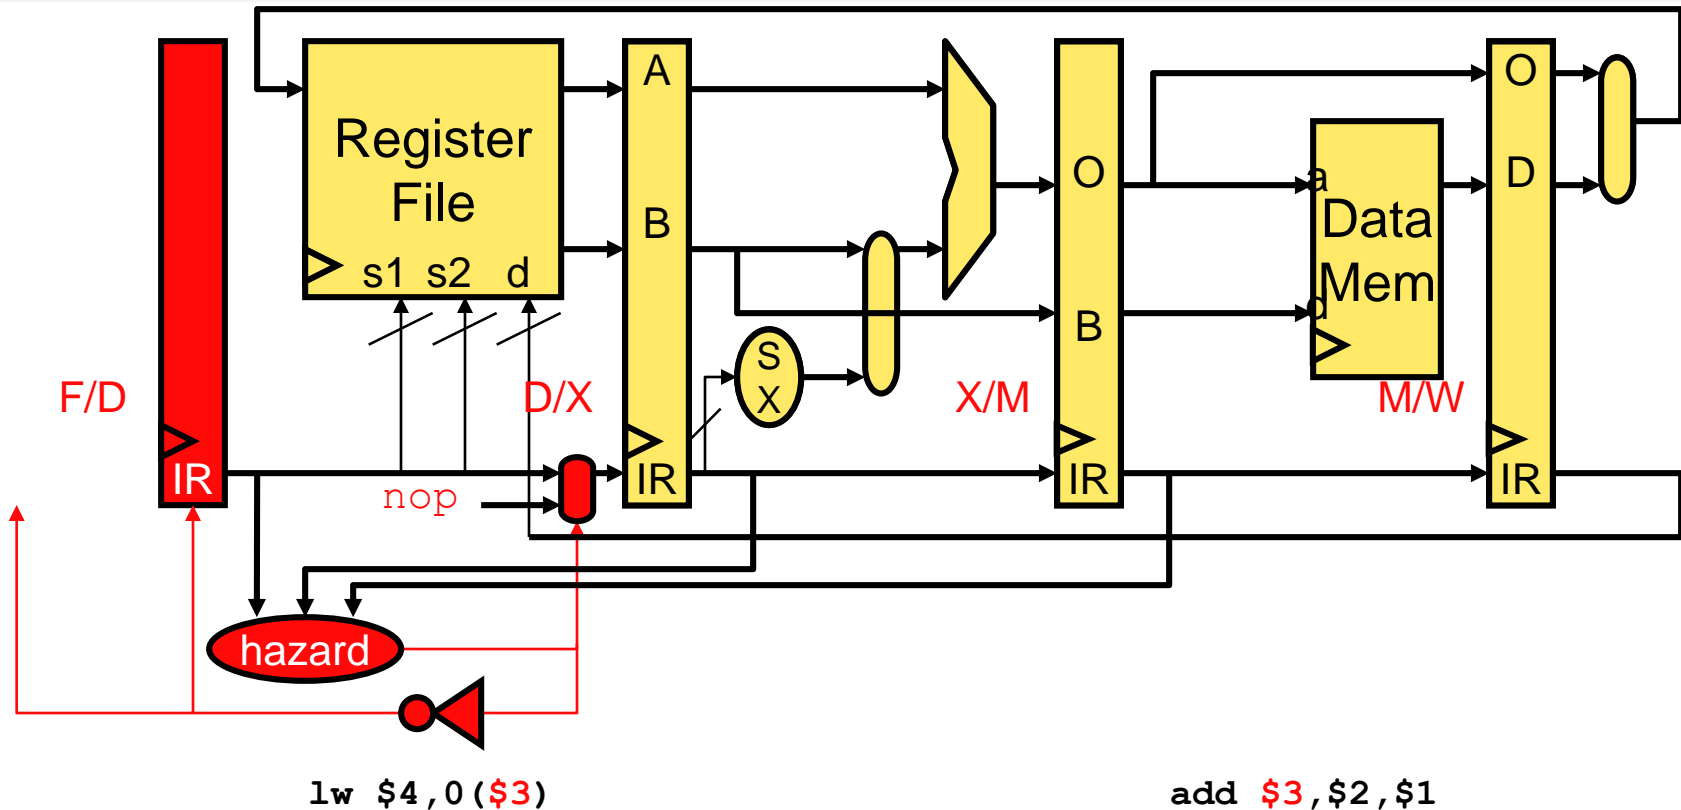
Hardware Interlock Example: cycle 1



$$\begin{aligned}
 & (\text{F/D.IR.RS1} == \text{D/X.IR.RD}) \parallel (\text{F/D.IR.RS2} == \text{D/X.IR.RD}) \parallel \\
 & (\text{F/D.IR.RS1} == \text{X/M.IR.RD}) \parallel (\text{F/D.IR.RS2} == \text{X/M.IR.RD})
 \end{aligned}$$

= 1

Hardware Interlock Example: cycle 2

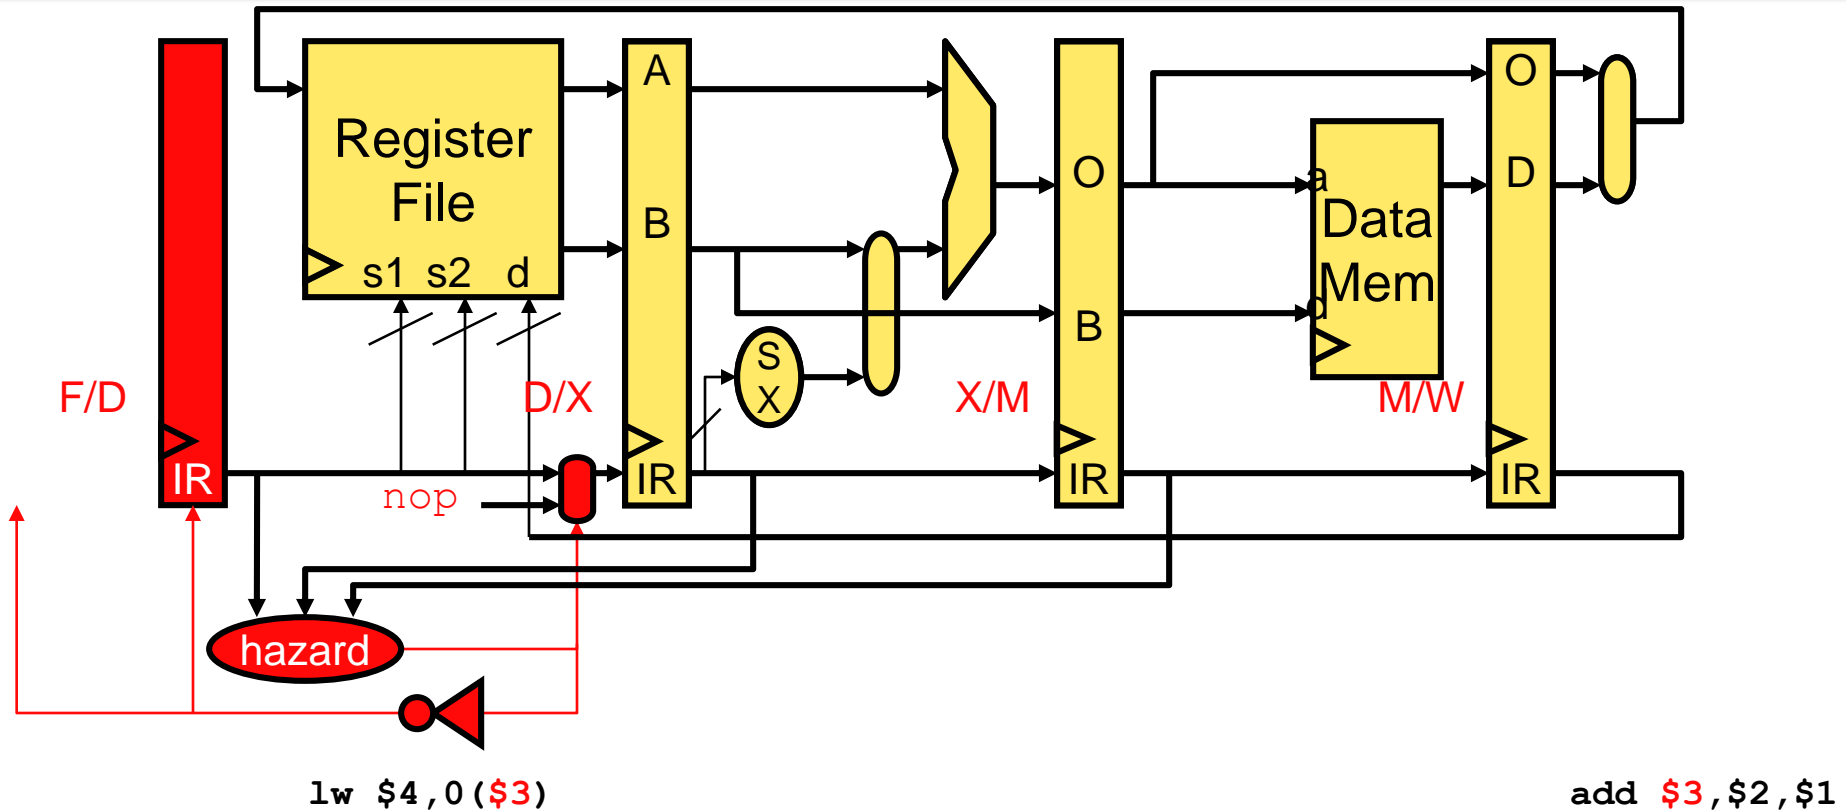


$$(F/D.IR.RS1 == D/X.IR.RD) \parallel (F/D.IR.RS2 == D/X.IR.RD) \parallel$$

$$(F/D.IR.RS1 == X/M.IR.RD) \parallel (F/D.IR.RS2 == X/M.IR.RD)$$

= 1

Hardware Interlock Example: cycle 3



$$(F/D.IR.RS1 == D/X.IR.RD) \parallel (F/D.IR.RS2 == D/X.IR.RD) \parallel (F/D.IR.RS1 == X/M.IR.RD) \parallel (F/D.IR.RS2 == X/M.IR.RD)$$

$$= 0$$

Pipeline Control Terminology

- Hardware interlock maneuver is called **stall** or **bubble**
- Mechanism is called **stall logic**
- Part of more general **pipeline control** mechanism
 - Controls advancement of insns through pipeline
- Distinguished from **pipelined datapath control**
 - Controls datapath at each stage
 - Pipeline control controls advancement of datapath control

Pipeline Diagram with Data Hazards

- Data hazard stall indicated with **d***
 - Stall propagates to younger insns

	1	2	3	4	5	6	7	8	9
add \$3,\$2,\$1	F	D	X	M	W				
lw \$4,0(\$3)		F	d*	d*	D	X	M	W	
sw \$6,4(\$7)					F	D	X	M	W


- This is not OK (why?)

	1	2	3	4	5	6	7	8	9
add \$3,\$2,\$1	F	D	X	M	W				
lw \$4,0(\$3)		F	d*	d*	D	X	M	W	
sw \$6,4(\$7)			F	D	X	M	W		

Hardware Interlock Performance


- Hardware interlocks: same as software interlocks
 - 20% of insns require 1 cycle stall (i.e., insertion of 1 `nop`)
 - 5% of insns require 2 cycle stall (i.e., insertion of 2 `nops`)
 - $\text{CPI} = 1 + 0.20 \times 1 + 0.05 \times 2 = \mathbf{1.3}$
 - So, either CPI stays at 1 and #insns increases 30% (software)
 - Or, #insns stays at 1 (relative) and CPI increases 30% (hardware)
 - Same difference
- Anyway, we can do better

Observe



- This situation seems broken
 - `lw $4, 0($3)` has already read `$3` from regfile
 - `add $3, $2, $1` hasn't yet written `$3` to regfile
- But fundamentally, everything is still OK
 - `lw $4, 0($3)` hasn't actually **used** `$3` yet
 - `add $3, $2, $1` has already **computed** `$3`

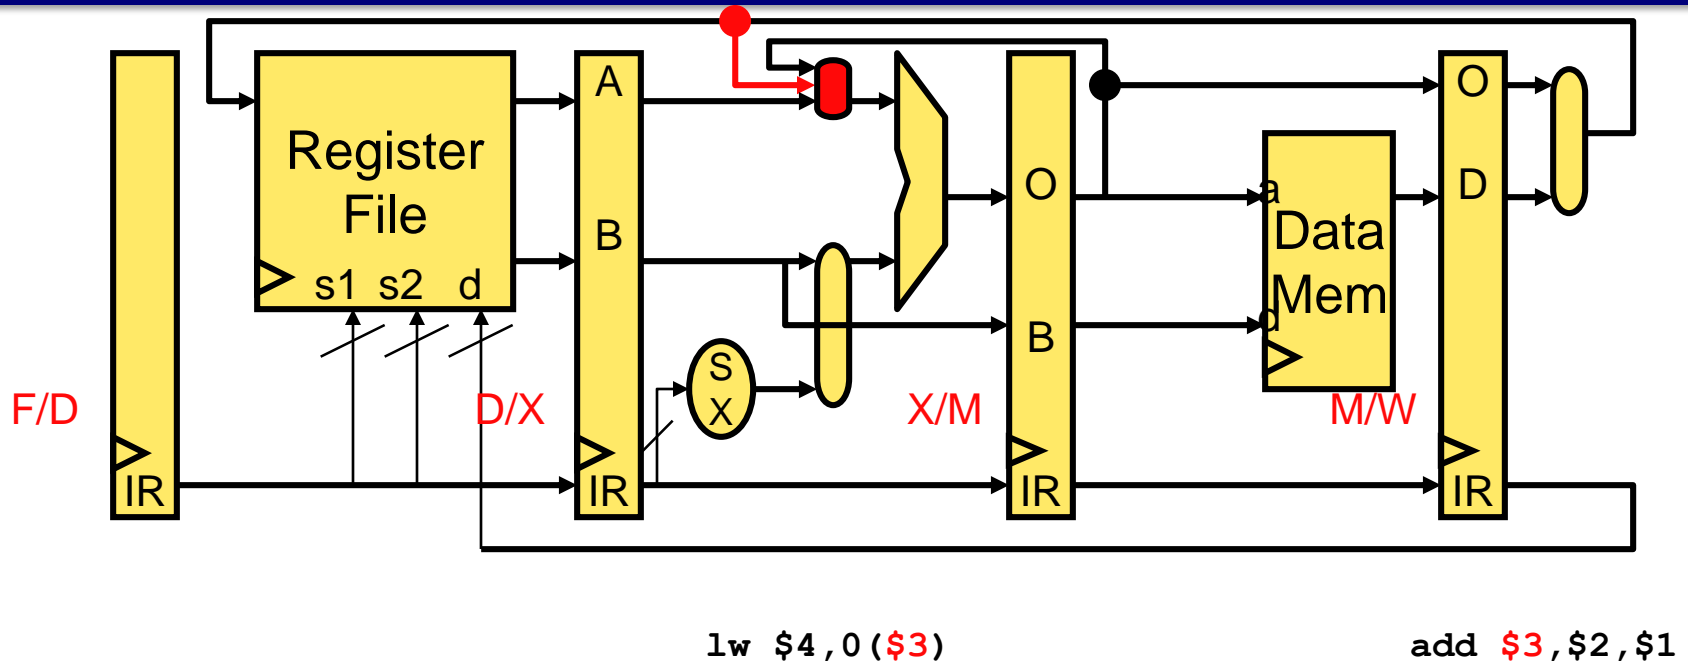
Bypassing



- **Bypassing**


- Reading a value from an intermediate (μ architectural) source
- Not waiting until it is available from primary source (RegFile)
- Here, we are bypassing the register file
- Also called **forwarding**

WX Bypassing




- What about this combination?
 - Add another bypass path and MUX input
 - First one was an **MX** bypass
 - This one is a **WX** bypass

ALUinB Bypassing



- Can also bypass to ALU input B

WM Bypassing?

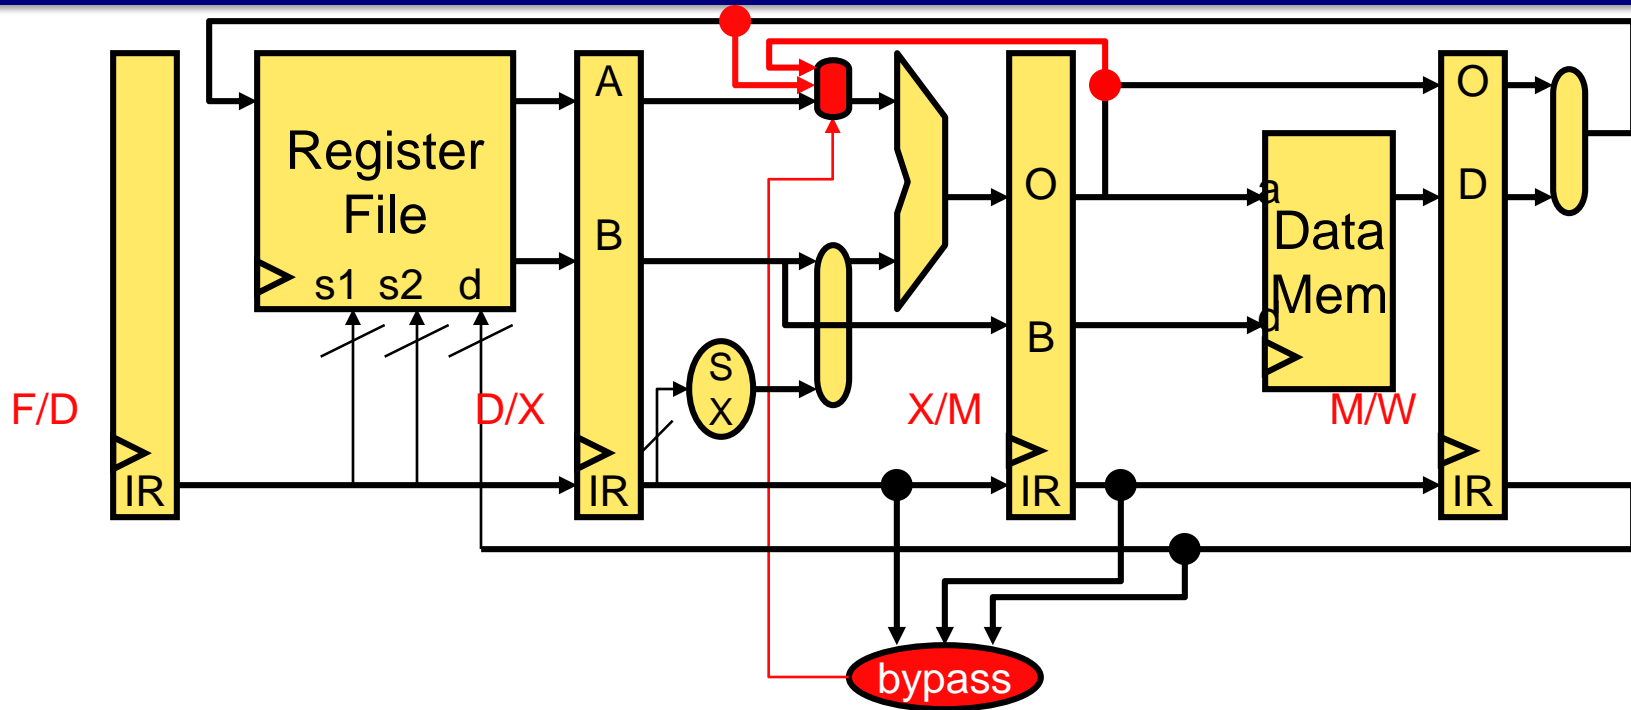


sw \$3, 0 (\$4)

lw \$3, 0 (\$2)

- Does WM bypassing make sense?
 - Yes, for the store data input

Bypass Logic

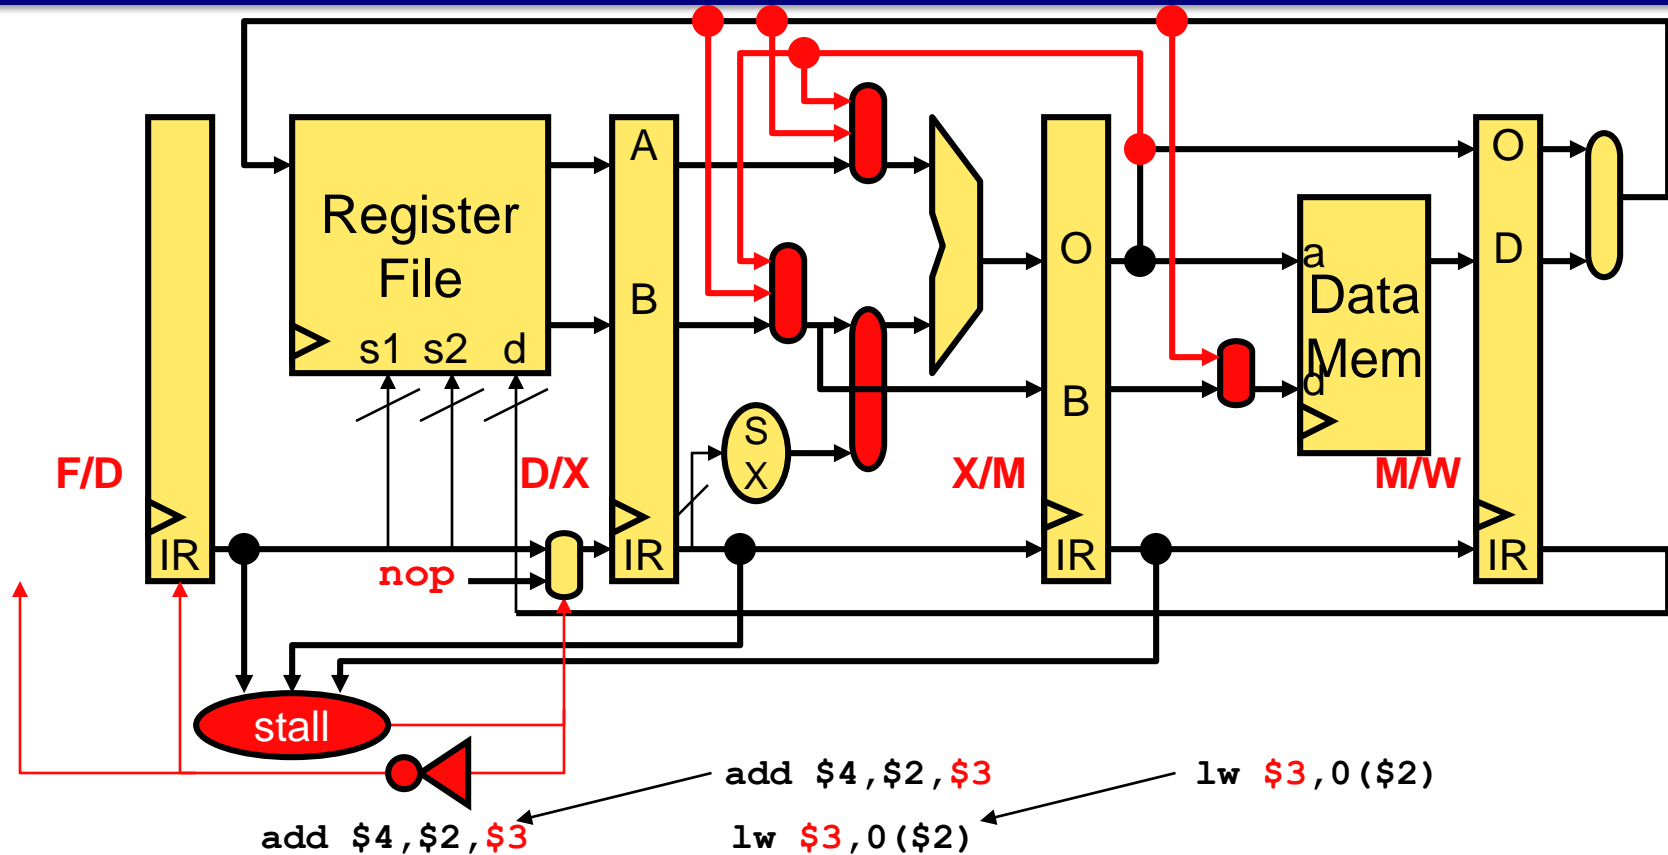


- Each MUX has its own control, here it is for MUX ALUinA
 $(D/X.IR.RS1 == X/M.IR.RD) \rightarrow \text{mux select} = 0$
 $(D/X.IR.RS1 == M/W.IR.RD) \rightarrow \text{mux select} = 1$
 Else $\rightarrow \text{mux select} = 2$

Bypass and Stall Logic

- Two separate things
 - Stall logic controls pipeline registers
 - Bypass logic controls muxes
- But complementary
 - For a given data hazard: if can't bypass, must stall
- **Full bypassing**: all bypasses possible
 - Is stall logic still necessary? Yes

Yes, Load Output to ALU Input



```

Stall = (D/X.IR.OP==LOAD) && (
    (F/D.IR.RS1==D/X.IR.RD) ||
    ((F/D.IR.RS2==D/X.IR.RD) && (F/D.IR.OP!=STORE))
)
    
```

Yes, Load Output to ALU Input

$$\text{Stall} = (\text{D/X.IR.OP} == \text{LOAD}) \ \&\& \ ($$

F/D
 $\text{add } \$4, \$3, \$2$

D/X
 $\text{lw } \$3, 0(\$2)$

$$\text{((F/D.IR.RS1} == \text{D/X.IR.RD))} \ || \ ||$$

$$\text{((F/D.IR.RS2} == \text{D/X.IR.RD))} \ \&\& \ (\text{F/D.IR.OP} != \text{STORE}))$$

$$)$$

$$\text{Stall} = (\text{D/X.IR.OP} == \text{LOAD}) \ \&\& \ ($$

F/D
 $\text{add } \$4, \$2, \$3$

D/X
 $\text{lw } \$3, 0(\$2)$

$$\text{(F/D.IR.RS1} == \text{D/X.IR.RD))} \ || \ ||$$

$$\text{((F/D.IR.RS2} == \text{D/X.IR.RD))} \ \&\& \ (\text{F/D.IR.OP} != \text{STORE}))$$

$$)$$

$$\text{Stall} = (\text{D/X.IR.OP} == \text{LOAD}) \ \&\& \ ($$

F/D
 $\text{sw } \$3, 0(\$4)$

D/X
 $\text{lw } \$3, 0(\$2)$

$$\text{(F/D.IR.RS1} == \text{D/X.IR.RD))} \ || \ ||$$

$$\text{((F/D.IR.RS2} == \text{D/X.IR.RD))} \ \&\& \ (\text{F/D.IR.OP} != \text{STORE}))$$

$$)$$

Intuition: "Stall if it's a load where rs1 is a data hazard for the next instruction, or where rs2 is a data hazard in a *non-store* next instruction". This is because rs2 is safe in a store instruction, because it doesn't use the X stage, and can be WM bypassed.

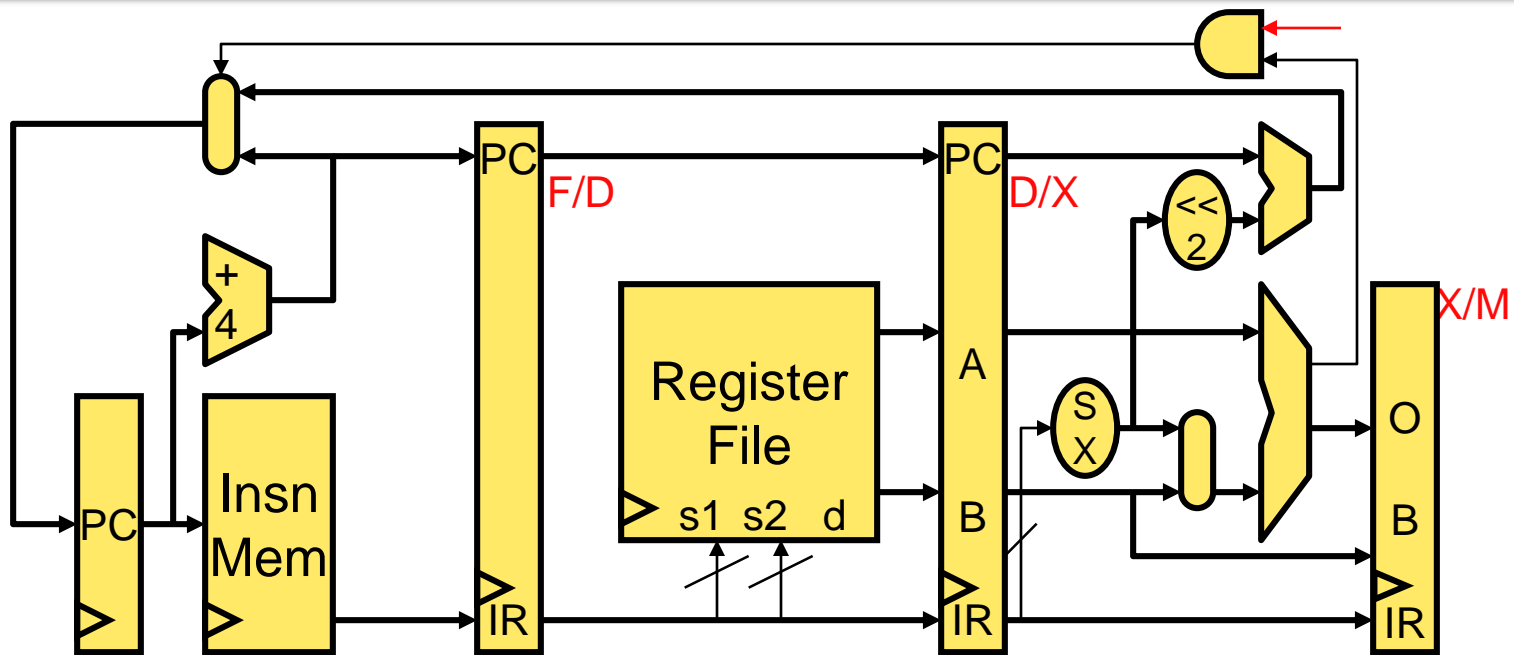
Pipeline Diagram With Bypassing

	1	2	3	4	5	6	7	8	9
add \$3,\$2,\$1	F	D	X	M	W				
lw \$4,0(\$3)		F	D	X	M	W			
addi \$6,\$4,1			F	d*	D	X	M	W	

- Sometimes you will see it like this
 - Denotes that stall logic implemented at X stage, rather than D
 - Equivalent, doesn't matter when you stall as long as you do

	1	2	3	4	5	6	7	8	9
add \$3,\$2,\$1	F	D	X	M	W				
lw \$4,0(\$3)		F	D	X	M	W			
addi \$6,\$4,1			F	D	d*	X	M	W	

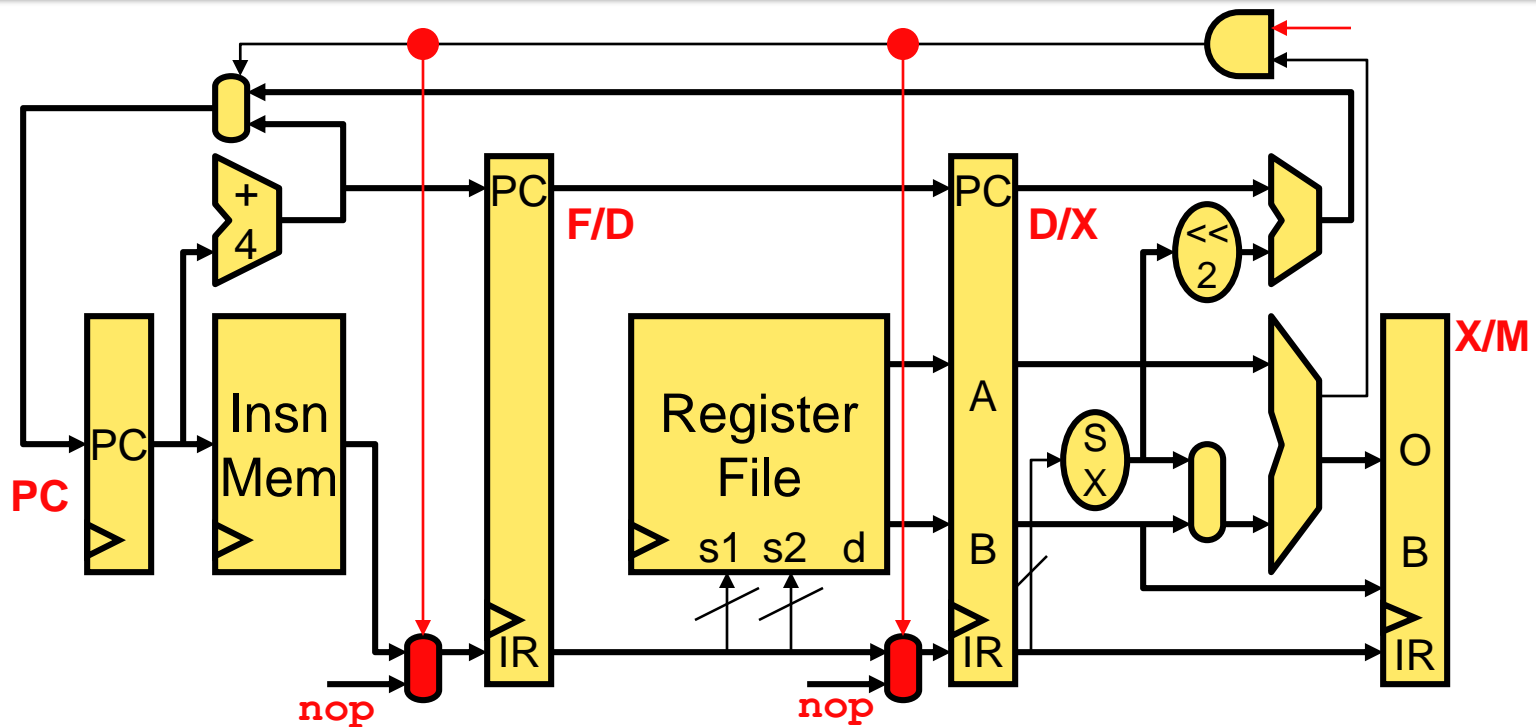
Control Hazards



- **Control hazards**

- Must fetch post branch insns before branch outcome is known
- Default: assume "**not-taken**" (at fetch, can't tell it's a branch)

Branch Recovery



- **Branch recovery:** what to do when branch **is** taken
 - **Flush** insns currently in F/D and D/X (they're wrong)
 - Replace with **NOPs**
 - + Haven't yet written to permanent state (RegFile, DMem)

Branch Recovery Pipeline Diagram

	1	2	3	4	5	6	7	8	9
addi \$3,\$0,1	F	D	X	M	W				
bnez \$3,targ		F	D	X	M	W			
sw \$6,4(\$7)			F	D					
addi \$8,\$7,1				F					
targ: sw \$6,4(\$7)					F	D	X	M	W

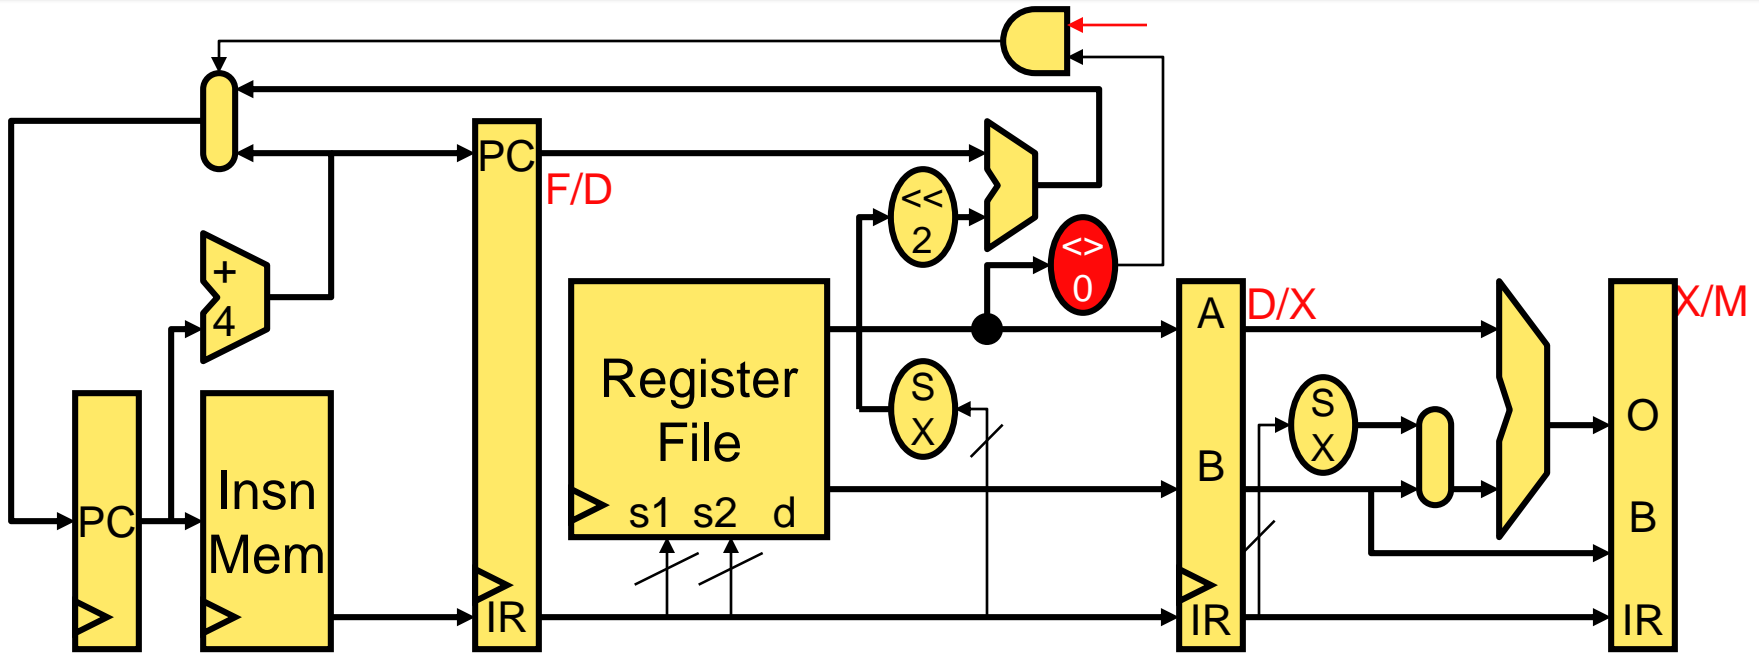
- Control hazards indicated with **c***
 - Penalty for taken branch is 2 cycles

	1	2	3	4	5	6	7	8	9
addi \$3,\$0,1	F	D	X	M	W				
bnez \$3,targ		F	D	X	M	W			
sw \$6,4(\$7)			c*	c*	F	D	X	M	W

Branch Performance

- Again, measure effect on CPI (clock period is fixed)
- Back of the envelope calculation
 - **Branch: 20%**, load: 20%, store: 10%, other: 50%
 - **75% of branches are taken**
- CPI if no branches = 1
- CPI with branches = $1 + 0.20 * 0.75 * 2 = 1.3$
 - **Branches cause 30% slowdown**
 - How do we reduce this penalty?

Fast Branch



- **Fast branch**: can decide at D, not X
 - Test must be comparison to zero or equality, no time for ALU
 - + New taken branch penalty is 1
 - Additional insns (**s1t**) for more complex tests
 - 25% of branches have complex tests that require extra insn
 - $CPI = 1 + 0.20 \cdot 0.75 \cdot \mathbf{1}(\text{branch}) + 0.20 \cdot 0.75 \cdot 0.25 \cdot 1(\text{extra insn}) = \mathbf{1.19}$

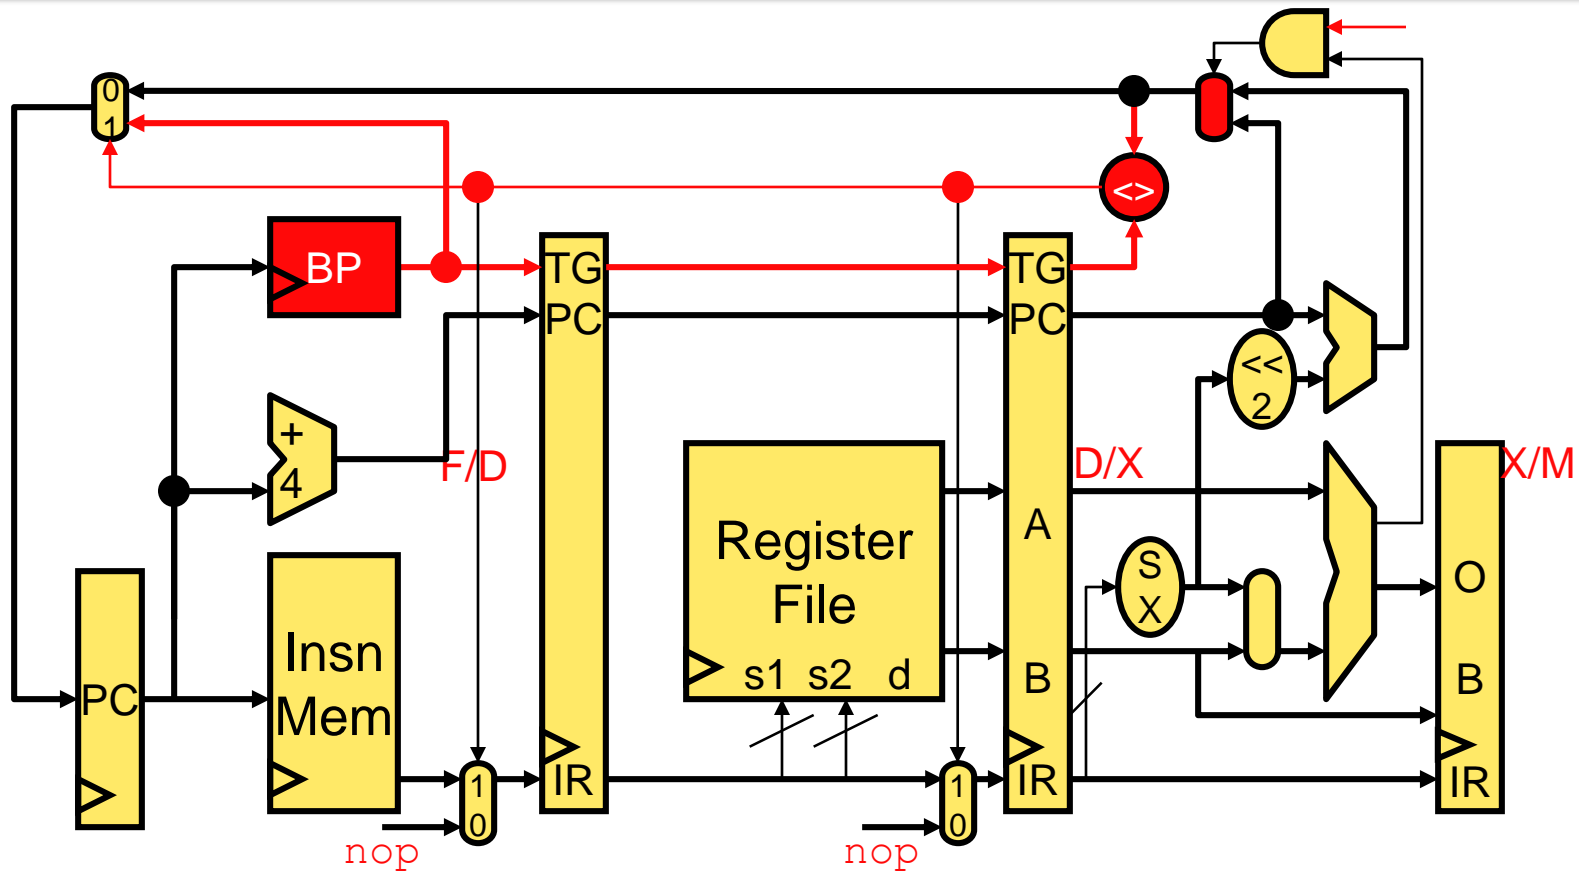
Speculative Execution

- Speculation: “risky transactions on chance of profit”
- **Speculative execution**
 - Execute before all parameters known with certainty
 - **Correct speculation**
 - + Avoid stall, improve performance
 - **Incorrect speculation (mis-speculation)**
 - Must abort/flush/squash incorrect insns
 - Must undo incorrect changes (recover pre-speculation state)
 - The “game”: $[\%_{\text{correct}} * \text{gain}] - [(1 - \%_{\text{correct}}) * \text{penalty}]$
- **Control speculation**: speculation aimed at control hazards

Control Speculation Mechanics

- Guess branch target, start fetching at guessed position
 - Doing nothing is implicitly guessing target is PC+4
 - Can actively guess other targets: **dynamic branch prediction**
- Execute branch to verify (check) guess
 - Correct speculation? keep going
 - Mis-speculation? Flush mis-speculated insns
 - Hopefully haven't modified permanent state (Regfile, DMem)
 - + Happens naturally in in-order 5-stage pipeline
- “Game” for in-order 5 stage pipeline
 - $\%_{\text{correct}} = ?$
 - Gain = 2 cycles
 - Penalty = 2 cycles

Dynamic Branch Prediction



- **Dynamic branch prediction:** guess outcome
 - Start fetching from guessed address
 - Flush on **mis-prediction** (notice new recovery circuit)

Branch Prediction: Short Summary

- Key principle of micro-architecture:
 - Programs do the same thing over and over (why?)
 - Exploit for performance:
 - Learn what a program did before
 - Guess that it will do the same thing again
- Inside a branch predictor: the short version
 - Use some of the PC bits as an **index** to a separate RAM
 - This RAM contains (a) branch destination and (b) whether we predict the branch will be taken
 - RAM is updated with results of past executions of branches
 - Algorithm for predictions can be simple ("assume it's same as last time"), or get quite fancy

Branch Prediction Performance

- Same parameters
 - **Branch: 20%**, load: 20%, store: 10%, other: 50%
 - 75% of branches are taken
- Dynamic branch prediction
 - Assume branches predicted with 75% accuracy (so 25% are penalized)
 - $\text{CPI} = 1 + 0.20 \times 0.25 \times 2 = \mathbf{1.1}$
- Branch (esp. direction) prediction was a hot research topic
 - Accuracies now 90-95%

Pipeline Depth

- No magic about 5 stages, trend had been to deeper pipelines
 - 486: 5 stages (50+ gate delays / clock)
 - Pentium: 7 stages
 - Pentium II/III: 12 stages
 - Pentium 4: 22 stages (~10 gate delays / clock) **“super-pipelining”**
 - Core1/2: 14 stages
- Increasing **pipeline depth**
 - + Increases clock frequency (reduces period)
 - But decreases IPC (increases CPI)
 - Branch mis-prediction penalty becomes longer
 - Non-bypassed data hazard stalls become longer
 - At some point, CPI losses offset clock gains, question is when?
 - 1GHz Pentium 4 was slower than 800 MHz PentiumIII
 - What was the point? People by frequency, not frequency * IPC

Real Pipelines...

- Real pipelines fancier than what we have seen
 - Superscalar: multiple instructions in a stage at once
 - Out-of-order: re-order instructions to reduce stalls
 - SMT: execute multiple threads at once on processor
 - Side by side, sharing pipeline resources
 - Multi-core: multiple pipelines on chip
 - Cache coherence (future lectures)

Summary

- Pipelining
 - Pipeline control