

# **ECE 550D**

## **Fundamentals of Computer Systems and Engineering**

### **Fall 2023**

## Memory Hierarchy

Xin Li & Dawei Liu  
Duke Kunshan University

Slides are derived from work by  
Andrew Hilton, Tyler Bletsch and Rabih Younes (Duke)

# Cache Miss Paper Simulation

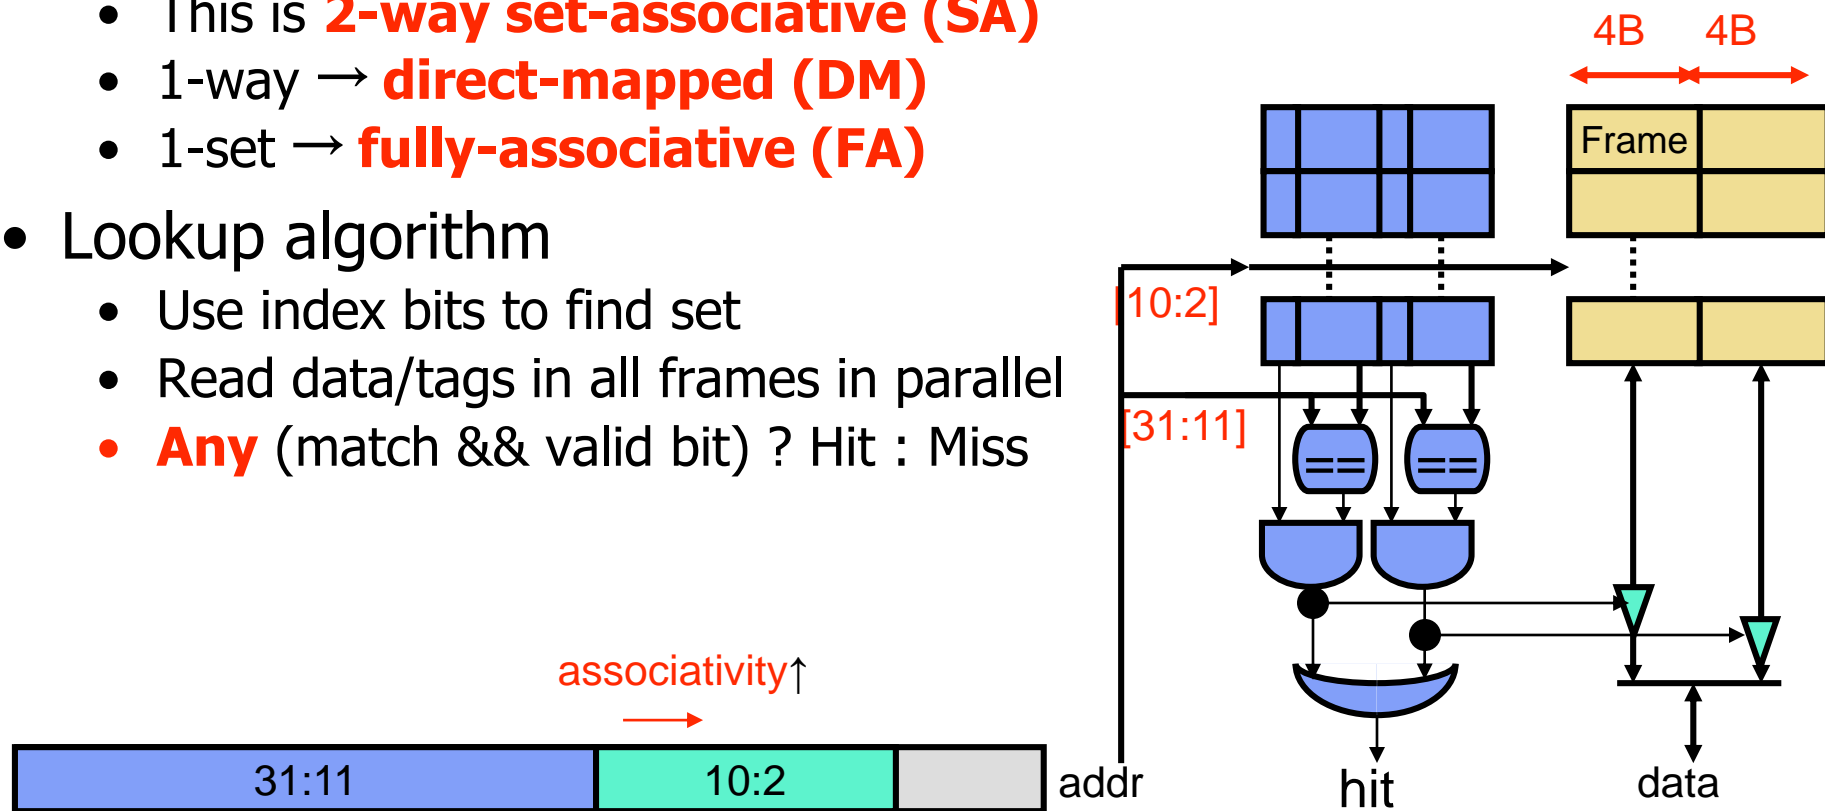
- 8B cache, 4B blocks -> 2 sets

Address	Tag	Index	Offset	Set 0	Set 1	Result
C	1	1	0	invalid	0	Miss
E	1	1	2	invalid	1	Hit
8	1	0	0	invalid	1	Miss
3	0	0	3	1	1	<b>Miss</b>
8	1	0	0	0	1	<b>Miss</b>
0	0	0	0	1	1	Miss
8	1	0	0	0	1	Miss
4	0	1	0	1	1	Miss
6	0	1	2	1	0	<b>Hit</b>

- 8 (1000) and 3 (0011): same set
- Can we do anything about this?

# Associativity

- New organizational dimension: **Associativity**
  - Block can reside in one of few frames
  - Frame groups called **sets**
  - Each frame in a set called a **way**
  - This is **2-way set-associative (SA)**
  - 1-way → **direct-mapped (DM)**
  - 1-set → **fully-associative (FA)**
- Lookup algorithm
  - Use index bits to find set
  - Read data/tags in all frames in parallel
  - **Any** (match && valid bit) ? Hit : Miss



# Cache Behavior 2-ways

Cache: 16b address, 4 sets, 2 ways, 4B blocks

Set #	Way 0			Way 1		
	V	Tag	Data	V	Tag	Data
0	0	000	00 00 00 00	0	000	00 00 00 00
1	0	000	00 00 00 00	0	000	00 00 00 00
2	0	000	00 00 00 00	0	000	00 00 00 00
3	0	000	00 00 00 00	0	000	00 00 00 00

# Cache Behavior 2-ways

Access address 0x1234

= 0001 0010 0011 0100

Tag = 123

Index = 1

Offset = 0

Set #	Way 0			Way 1		
	V	Tag	Data	V	Tag	Data
0	0	000	00 00 00 00	0	000	00 00 00 00
1	1	123	00 00 00 00	0	000	00 00 00 00
2	0	000	00 00 00 00	0	000	00 00 00 00
3	0	000	00 00 00 00	0	000	00 00 00 00

Miss. Request from next level. Wait...

# Cache Behavior 2-ways

Access address 0x2234

= 0001 0010 0011 0100

Tag = 223

Index = 1

Offset = 0

Set #	Way 0			Way 1		
	V	Tag	Data	V	Tag	Data
0	0	000	00 00 00 00	0	000	00 00 00 00
1	1	123	0F 1E 39 EC	1	223	00 00 00 00
2	0	000	00 00 00 00	0	000	00 00 00 00
3	0	000	00 00 00 00	0	000	00 00 00 00

Miss. Request from next level. Wait...

# Cache Behavior 2-ways

Access address 0x1234

= 0001 0010 0011 0100

Tag = 123

Index = 1

Offset = 0

Set #	Way 0			Way 1		
	V	Tag	Data	V	Tag	Data
0	0	000	00 00 00 00	0	000	00 00 00 00
1	1	123	0F 1E 39 EC	1	223	01 CF D0 87
2	0	000	00 00 00 00	0	000	00 00 00 00
3	0	000	00 00 00 00	0	000	00 00 00 00

Hit. In Way 0

# Cache Miss Paper Simulation

- 4b address, 8B cache, 2B blocks, 2 ways, 2 sets

Offset:  $\log_2 2 = 1$  bit

Index:  $\log_2 2 = 1$  bit

Tag:  $4 - 1 - 1 = 2$  bit

- What happens for each request?



# Cache Miss Paper Simulation

- 4b address, 8B cache, 2B blocks, 2 ways, 2 sets

				Set 0		Set 1		
Address	Tag	Index	Offset	Way0	Way1	Way0	Way1	Result
C	3	0	0	invalid	0	0	1	Miss
E	3	1	0	3	0	0	1	Miss
8	2	0	0	3	0	0	3	Miss
3	0	1	1	3	2	0	3	Hit
8	2	0	0	3	2	0	3	Hit
0	0	0	0	3	2	0	3	Miss
8	2	0	0	0	2	0	3	Hit
4	1	0	0	0	2	0	3	Miss
6	1	1	0	1	2	0	3	Miss

- What happens for each request?

# Cache structure math summary


- Given capacity, block\_size, ways (associativity), and word\_size.
- Cache parameters:
  - $\text{num\_frames} = \text{capacity} / \text{block\_size}$
  - $\text{sets} = \text{num\_frames} / \text{ways} = \text{capacity} / \text{block\_size} / \text{ways}$
- Address bit fields:
  - $\text{offset\_bits} = \log_2(\text{block\_size})$
  - $\text{index\_bits} = \log_2(\text{sets})$
  - $\text{tag\_bits} = \text{word\_size} - \text{index\_bits} - \text{offset\_bits}$
- Numeric way to get offset/index/tag from address:
  - $\text{block\_offset} = \text{addr} \% \text{block\_size}$
  - $\text{index} = (\text{addr} / \text{block\_size}) \% \text{sets}$
  - $\text{tag} = \text{addr} / (\text{sets} * \text{block\_size})$

# Replacement Policies

- Set-associative caches present a new design choice
  - On cache miss, which block in set to replace (kick out)?
- Belady's (oracle): block that will be used furthest in future
- Random
- FIFO (first-in first-out)
- LRU (least recently used)
  - Fits with temporal locality, LRU = least likely to be used in future
- **NMRU (not most recently used)**
  - An easier to implement approximation of LRU
  - Equal to LRU for 2-way SA caches


# NMRU Implementation

- Add **MRU** field to each set
  - MRU data is encoded "way"
  - Hit? update MRU
  - Fill? write enable  $\sim$ MRU



# Associativity And Performance

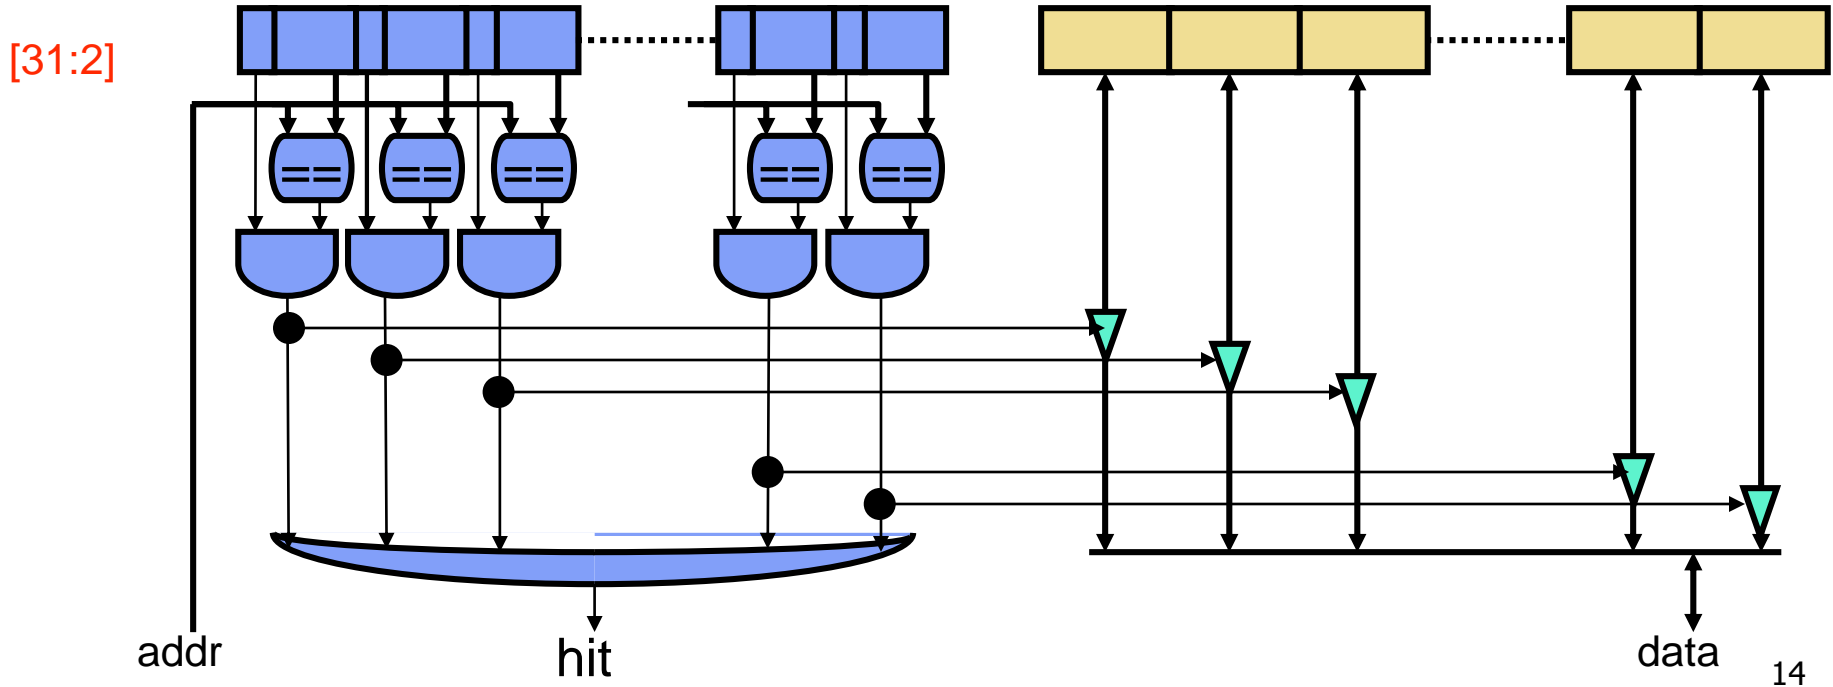
- The associativity game
  - + Higher associative caches have lower  $\%_{\text{miss}}$
  - $t_{\text{hit}}$  increases
    - Additional logic at output
    - But not much for low associativities (2,3,4,5)
  - $t_{\text{avg}}$ ?



- Block-size and number of sets should be powers of two
  - Makes indexing easier (just rip bits out of the address)
- 5-way set-associativity? No problem


# Full Associativity

- How to implement full (or at least high) associativity?
  - This way is terribly inefficient
  - 1K matches are unavoidable, but 1K data reads + 1K-to-1 mux?




# Full-Associativity with CAMs


- **CAM**: content-addressable memory
  - Array of words with built-in comparators
  - Input is data (tag)
  - Output is 1H encoding of matching slot
- Fully associative cache
  - Tags as CAM, data as RAM
  - Effective but expensive (EE reasons)
  - Upshot: used for 16-/32-way associativity
    - No good way to build 1024-way associativity
    - + No real need for it, either



no index bits




# CAM -> Content Addressable Memory



- Input: Data to match
  - (example on left: 3 bits)
- Output: matching entries
  - (example on left: 4 entries)




# CAM circuit



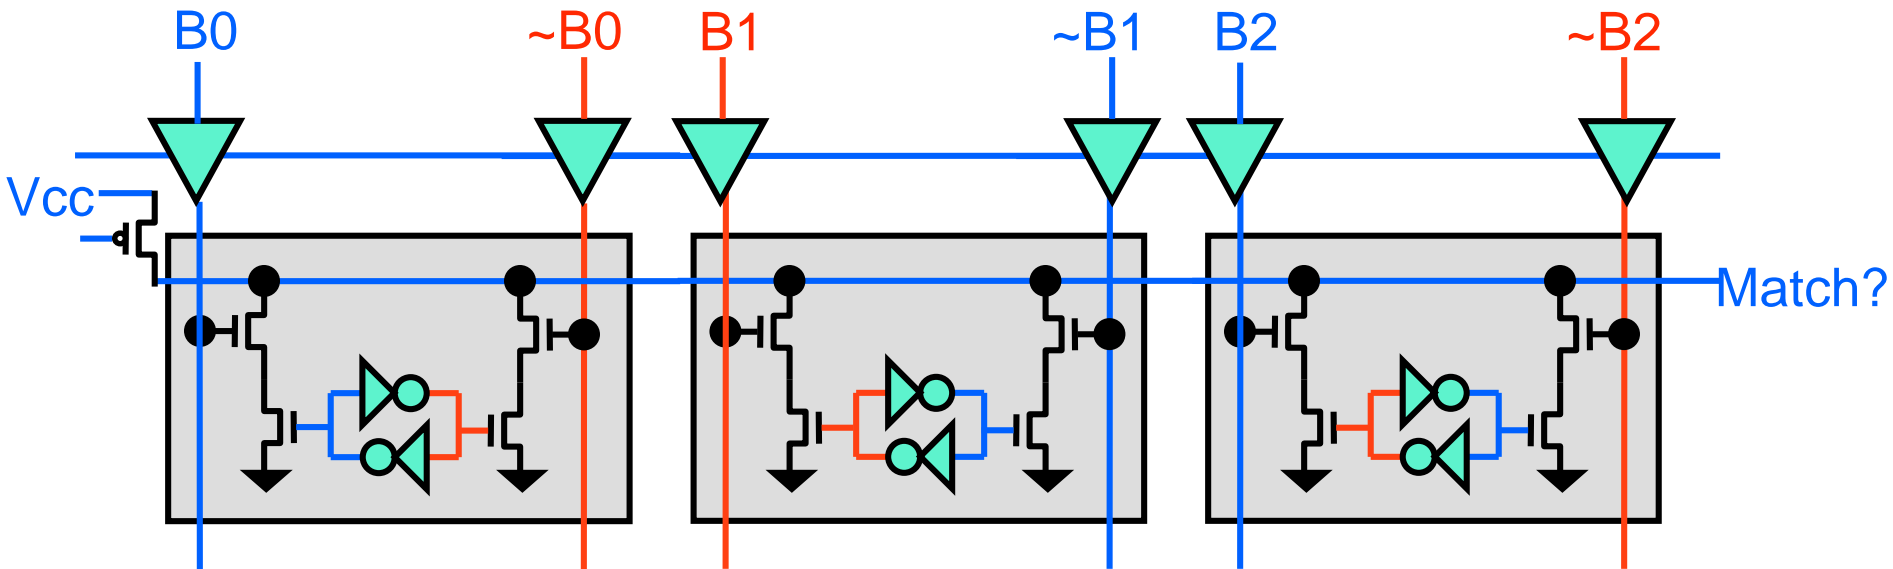
- CAM match port looks different from RAM r/w port
- Cells look similar
  - Note: Bit stored on right, ~Bit on left (opposite of inputs)

# CAM circuit



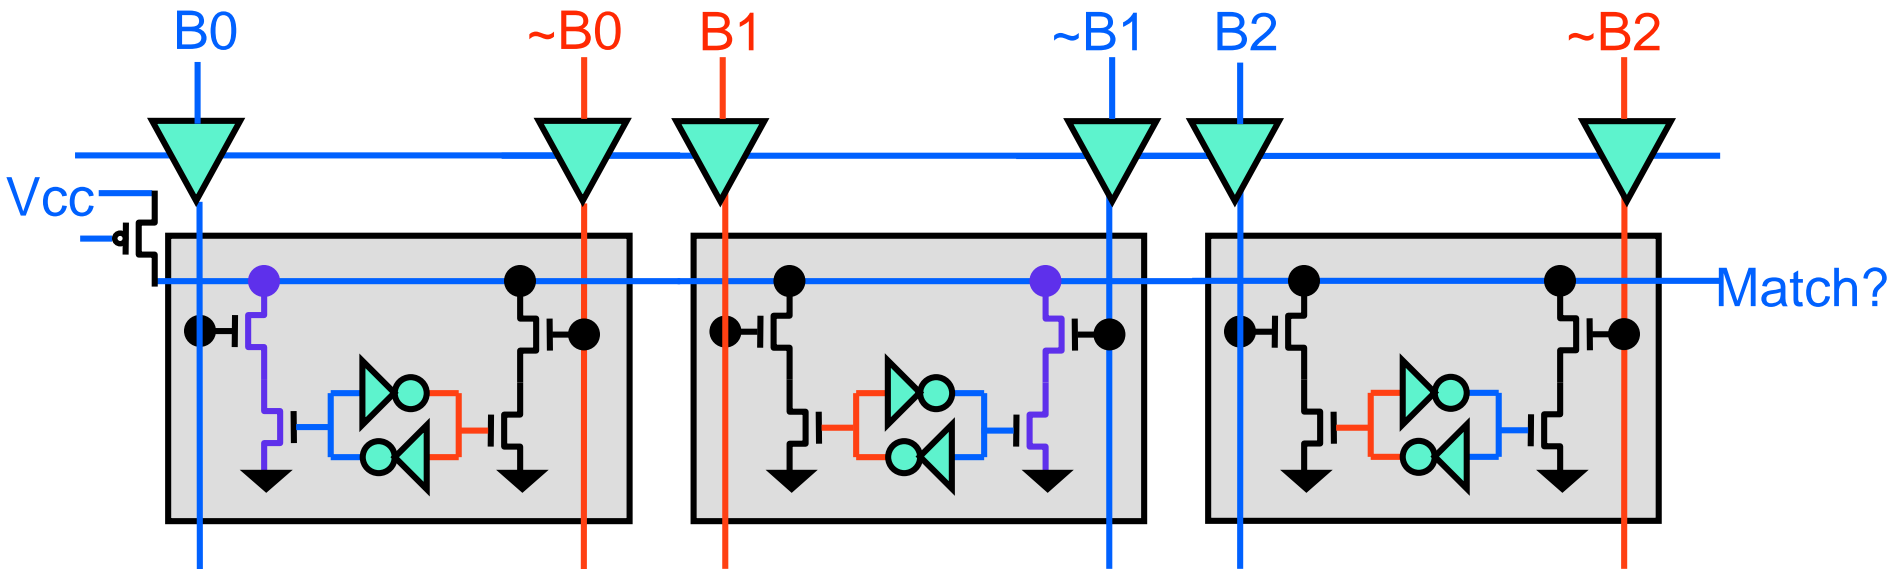
- CAM match port looks different from RAM r/w port
- Cells look similar
  - Note: Bit stored on right, ~Bit on left (opposite of inputs)
- Step 1: Precharge match line to 1 (first half of cycle)

# CAM circuit



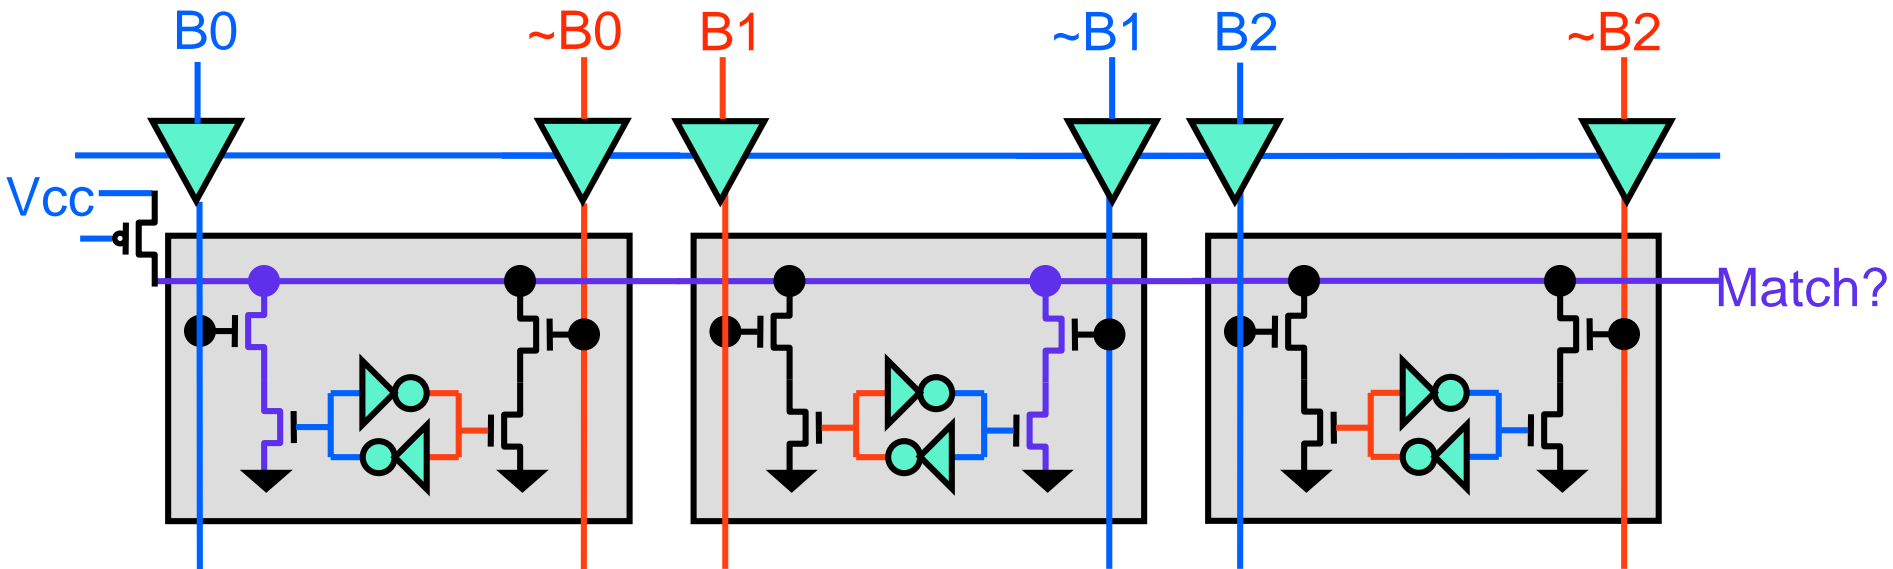
- CAM match port looks different from RAM r/w port
- Cells look similar
  - Note: Bit stored on right, ~Bit on left (opposite of inputs)
- Step 1: Precharge match line to 1 (first half of cycle)
- Step 2: Send data/~data down bit lines

# CAM circuit



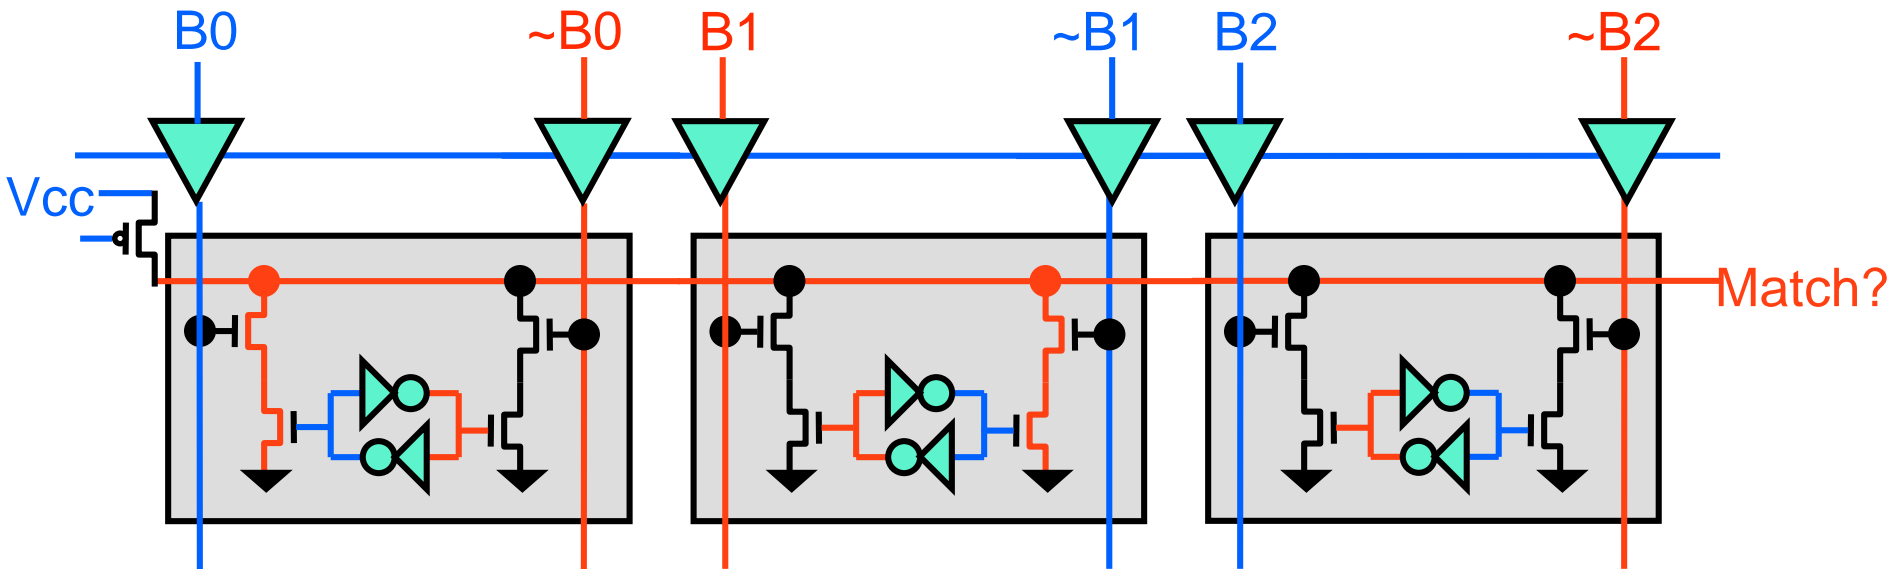
- CAM match port looks different from RAM r/w port
- Cells look similar
  - Note: Bit stored on right, ~Bit on left (opposite of inputs)
- Step 1: Precharge match line to 1 (first half of cycle)
- Step 2: Send data/~data down bit lines
  - Two 1s on same side (bit line  $\neq$  data) open NMOS path  $\rightarrow$  gnd

# CAM circuit




- CAM match port looks different from RAM r/w port
- Cells look similar
  - Note: Bit stored on right, ~Bit on left (opposite of inputs)
- Step 1: Precharge match line to 1 (first half of cycle)
- Step 2: Send data/~data down bit lines
  - Two 1s on same side (bit line  $\neq$  data) open NMOS path  $\rightarrow$  gnd
    - Drains match line 1  $\rightarrow$  0

# CAM circuit



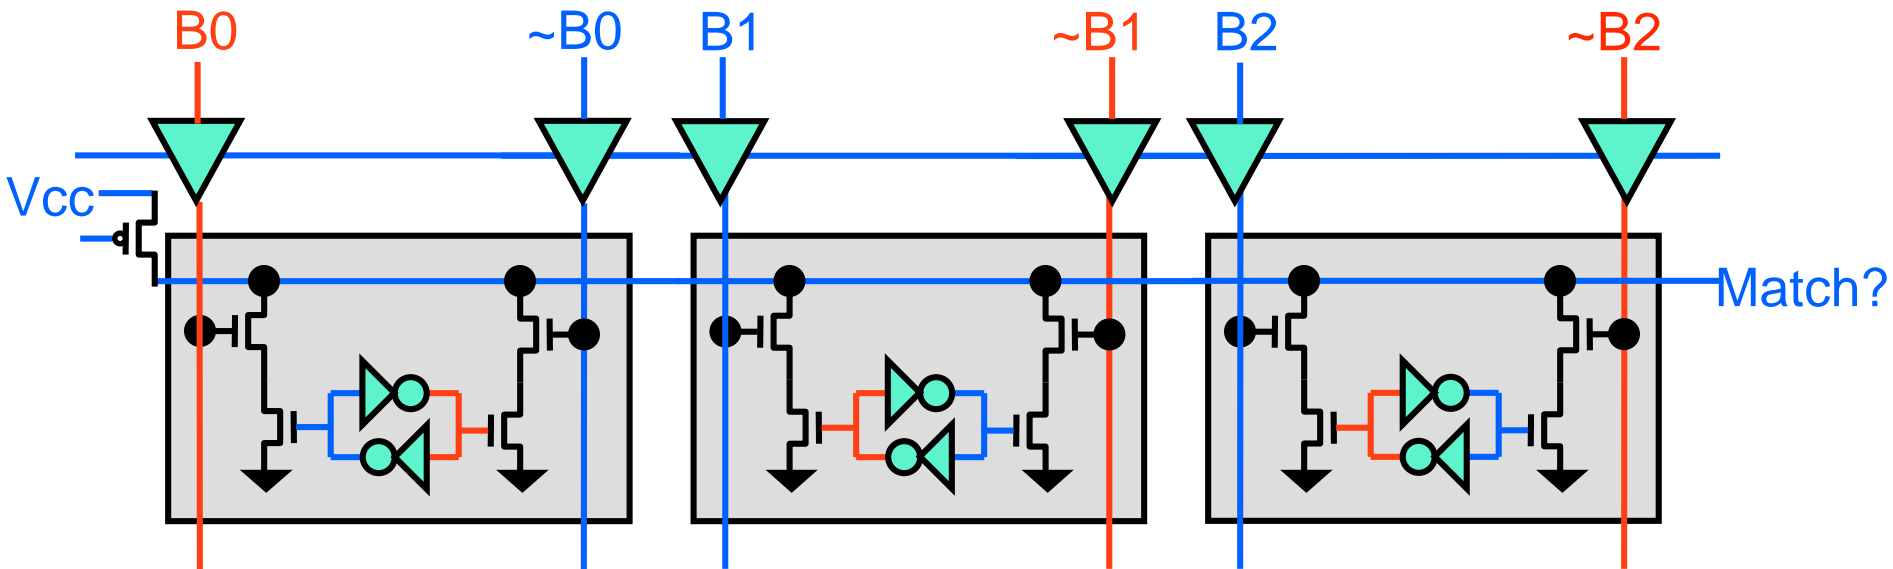
- CAM match port looks different from RAM r/w port
- Cells look similar
  - Note: Bit stored on right, ~Bit on left (opposite of inputs)
- Step 1: Precharge match line to 1 (first half of cycle)
- Step 2: Send data/~data down bit lines
  - Two 1s on same side (bit line  $\neq$  data) open NMOS path  $\rightarrow$  gnd
    - Drains match line 1  $\rightarrow$  0

# CAM circuit



- Note that if all bits match, each side has a 1 and a 0
  - One NMOS in the path from Match -> Gnd is closed
  - No conductive path -> Match keeps its charge @ 1

# CAMs: Slow and High Power..



- CAMs are slow and high power
  - Pre-charge all, discharge most match lines every search
    - Pre-charge + discharge take time: capacitive load of match line
  - Bit lines have high capacitive load: Driving 1 transistor per row



# ABC

- **Capacity**
  - + Decreases **capacity** misses
  - Increases  $t_{\text{hit}}$
- **Associativity**
  - + Decreases **conflict** misses
  - Increases  $t_{\text{hit}}$
- **Block size**
  - Increases **conflict** misses
  - + Decreases **compulsory** misses
  - ± Increases or decreases **capacity** misses
    - Little effect on  $t_{\text{hit}}$ , may exacerbate  $t_{\text{miss}}$
- How much they help depends...

*What are all these **blue** words? We'll see...*

# Different Problems -> Different Solutions

- Suppose we have 16B, direct-mapped cache w/ 4B blocks
  - 4 sets
  - Examine some access patterns and think about what would help
  - Misses in red
- Access pattern A:
  - As is: 0, 2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26
  - 8B blocks? 0, 2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26
  - 2-way assoc? 0, 2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26
- Access pattern B:
  - As is: 0, 128, 1, 129, 2, 130, 3, 131, 4, 132, 5, 133, 6
  - 8B blocks? 0, 128, 1, 129, 2, 130, 3, 131, 4, 132, 5, 133, 6
  - 2-way assoc? 0, 128, 1, 129, 2, 130, 3, 131, 4, 132, 5, 133, 6
- Access pattern C:
  - 0, 20, 40, 60, 48, 36, 24, 12, 1, 21, 41, 61, 49, 37, 25, 13, 2, 22, 42, 62, 50, 38, ...

# Analyzing Misses: 3C Model (Hill)

- Divide cache misses into categories based on cause
  - **Compulsory**: block size is too small (i.e., address not seen before)
  - **Capacity**: capacity is too small
  - **Conflict**: associativity is too low

# Different Problems -> Different Solutions

- Access pattern A: Compulsory misses
  - 0, 2, 4, 6, 8, 10, 12, 14, 16, 18, 20, 22, 24, 26
  - For misses, have not accessed that block
    - Capacity/associativity won't help (never had it)
    - Larger block -> include more data in one block -> more hits
- Recognizing compulsory misses
  - Never seen the **block** before

# Different Problems -> Different Solutions

- Access pattern B: Conflict misses
  - 0, 128, 1, 129, 2, 130, 3, 131, 4, 132, 5, 133, 6
  - 0 and 128 map to same set (set 0): kick each other out ("conflict")
  - Larger block? No help
  - Larger cache? Only helps if MUCH larger (256B instead of 16B)
  - Higher associativity? Fixes problem
    - Can have both 0 and 128 in set 0 at same time (different ways)
- Recognizing conflict misses:
  - Count unique blocks between last access and miss (inclusive)
  - Number of unique blocks  $\leq$  number of blocks in cache? Conflict
    - Enough space to hold them all...
    - Just must be having set conflict

# Different Problems -> Different Solutions

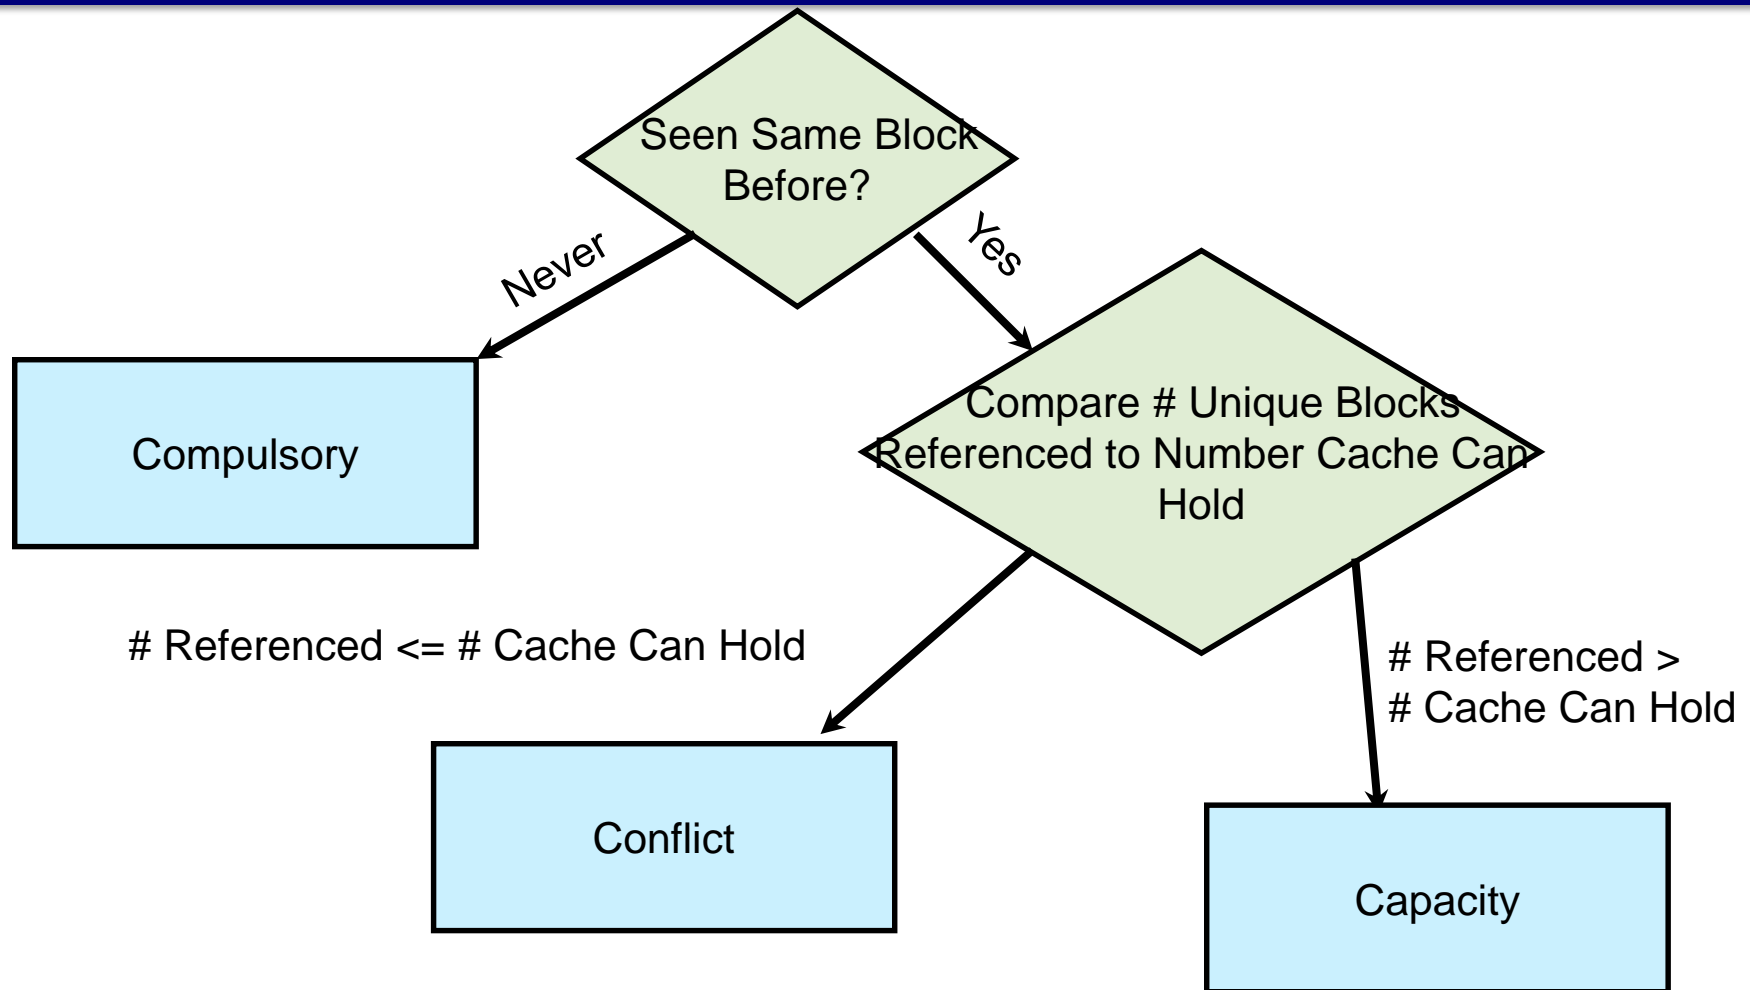
- Access pattern C: Capacity Misses

- 0,20,40,60,48,36,24,12,1,21,41,61,49,37,25,13,2,22,42,62,50,38,...
- Larger block size? No help
  - Even 16B block (entire cache) won't help
- Associativity? No help... even at full assoc
  - After 0, 20, 40, 60: kick out 0 for 48
  - Kick out 20 for 36
  - Kick out 40 for 24...
- Solution: make cache larger
  - Doubling cache size turns all most misses into hits
    - A few compulsory misses remain
- 0,20,40,60,48,36,24,12,1,21,41,61,49,37,25,13,2,22,42,62,50,38,...

- Recognizing Capacity Misses

- Count unique blocks between last access and miss (inclusive)
- Number of unique blocks > number of blocks in cache? Capacity
  - Just can't hold them all

# Miss Categorization Flow Chart



- **Capacity**
  - + Decreases capacity misses
  - Increases  $t_{\text{hit}}$
- **Associativity**
  - + Decreases conflict misses
  - Increases  $t_{\text{hit}}$
- **Block size**
  - Increases conflict misses
  - + Decreases compulsory misses
  - ± Increases or decreases capacity misses
    - Little effect on  $t_{\text{hit}}$ , may exacerbate  $t_{\text{miss}}$
- How much they help depends...




# Two Optimizations

- **Victim buffer**: for conflict misses
  - Technically: reduces  $t_{\text{miss}}$  for these misses, doesn't eliminate them
  - Depends how you do your accounting
- **Prefetching**: for capacity/compulsory misses

# Victim Buffer

- Conflict misses: not enough associativity
  - High associativity is expensive, but also rarely needed
- **Victim buffer (VB)**: small FA cache (e.g., 4 entries)
  - Small so very fast
  - Blocks kicked out of cache placed in VB
  - On miss, check VB: hit ? Place block back in cache
  - 4 extra ways, shared among all sets
    - + Only a few sets will need it at any given time
  - On cache fill path: reduces  $t_{\text{miss}}$ , no impact on  $t_{\text{hit}}$
  - + Very effective in practice



# Prefetching

- **Prefetching**: put blocks in cache proactively/speculatively
  - In software: insert prefetch (non-binding load) insns into code
  - In hardware: cache controller generates prefetch addresses
- Keys: anticipate upcoming miss addresses accurately
  - **Timeliness**: initiate prefetches sufficiently in advance
  - **Accuracy**: don't evict useful data
  - Prioritize misses over prefetches
- Simple algorithm: **next block prefetching**
  - Miss address **X** → prefetch address **X+block\_size**
  - Works for insns: sequential execution
    - What about non-sequential execution?
  - Works for data: arrays
    - What about other data-structures?
  - Address prediction is actively researched area

